

GNU/Linux assembly language

Reviewing the basics of assembly
language programming for Intel
x86-based Linux systems

Our motivations

- We want to study the new Intel processor technologies for Core-2 Duo & Pentium-D:
 - EM64T (Extended Memory 64-bit technology)
 - VT (Virtualization Technology)
- These capabilities are optionally ‘enabled’ during the system’s ‘startup’ phase, so we will want to execute our own ‘boot loader’ code to have control at the earliest stage

‘boot’ code

- On PC systems the mechanism for doing the IPL (‘Initial Program Load’) requires a ‘boot-loader’ program that can fit within a single 512-byte disk-sector
- High-level programming languages (such C/C++) typically employ compilers which generate object-code files whose size is too large to be stored in a disk-sector


Portability implications

- Also 'High-Level' languages typically have platform 'portability' as a key design-goal, but they can't achieve that requirement if architecture-specific language-constructs are employed; thus they usually offer to a systems programmer only those features in the 'lowest common denominator' of the various different computer-systems

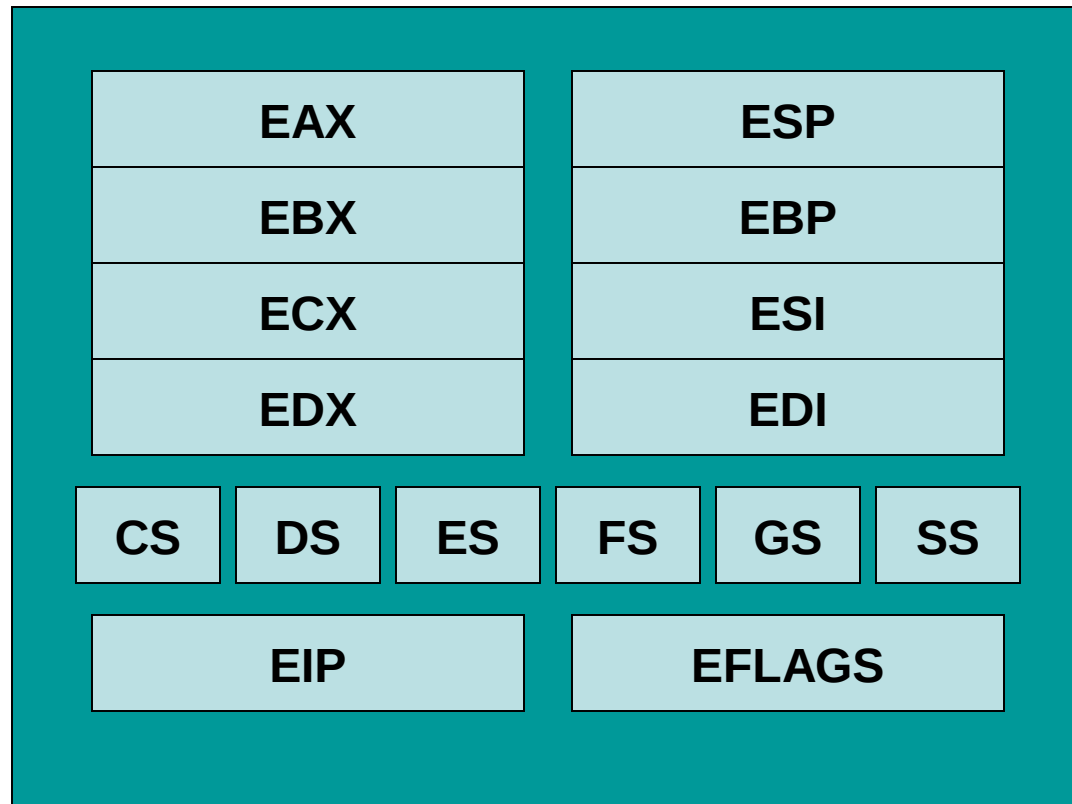
Intel's VMX instructions

- So when we want to explore Intel's VM-x technology (Virtualization for x86 CPUs), we will need to use 'low-level' computer language (i.e., assembly language)
- The advantage, of course, is that we will thereby acquire total control over the CPU in any Intel-based computer that supports the new so-called VMX instruction-set.

Simplified Block Diagram




The sixteen x86 registers



Intel Pentium processor

The x86 EFLAGS register



Legend:

IOPL = I/O Privilege-Level (0,1,2,3)

AC = Alignment-Check (1=yes, 0=no)

NT = Nested-Task (1=yes, 0=no)

RF = Resume Flag (1=yes, 0=no)

VM = Virtual-8086 Mode (1=yes, 0=no)

VIF = 'Virtual' Interrupt-Flag VIP = 'Virtual' Interrupt is Pending

ID = the CPUID-instruction is implemented if this bit can be 'toggled'

ZF = Zero Flag

SF = Sign Flag

CF = Carry Flag

PF = Parity Flag

OF = Overflow Flag

AF = Auxiliary Flag

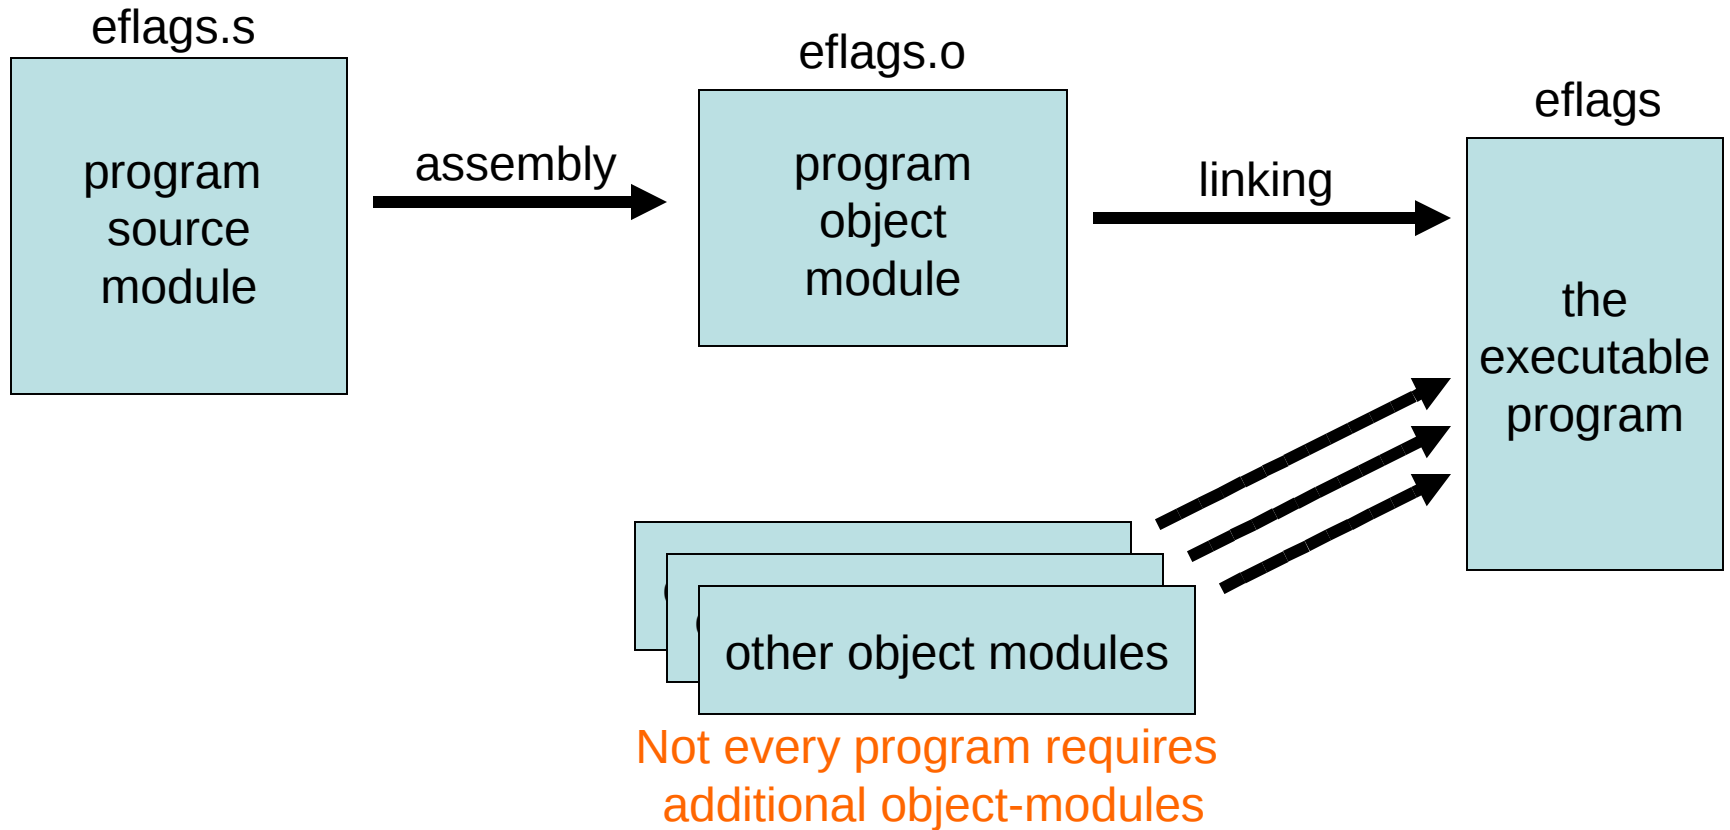
 = 'reserved' bit

 = 'status' flag (for application programming)

Our 'eflags.s' demo

- This program shows the EFLAGS bits at the moment the program began executing
- It's instructive to run this program before you execute our 'iopl3' command, and do so again after you execute that command

program translation steps



program translation commands

- To 'assemble' your source-file:


```
$ as eflags.s -o eflags.o
```

- To 'link' the resulting object-file:

```
$ ld eflags.o -o eflags
```

- To 'run' your executable program:

```
$ ./eflags
```



Last night (Wed, 24 Jan 2007) Alex Fedosov wrote:

>

> Subject: Re: hardware

>

> I checked #11-#12 serial cable and it seems that it was

> loose. Let me know if it continues to be a problem.

>

> Good news! Your machines are here! We will be building

> them tomorrow.

>

> -a

Our 'feedback.s' demo

- We wrote an example-program that shows how you could program the serial UART in the GNU/Linux assembly language
- It uses the x86 'in' and 'out' instructions to access the UART's i/o-ports:
 - 0x03F8: port-address for Divisor-Latch
 - 0x03FB: port-address for Line-Control
 - 0x03FD: port-address for Line-Status
 - 0x03F8: port-address for RxD and TxD

Statement layout

| | | | |
|---------------|---------------|-------------------|------------------|
| label: | opcode | operand(s) | # comment |
|---------------|---------------|-------------------|------------------|

Parsing an assembly language statement

- A colon character (':') terminates the label
- A white-space character (e.g., blank or tab) separates the opcode from its operand(s)
- A hash character ('#') begins the comment

How 'outb' works

- Three-steps to perform 'outb(data, port);'
 - Step 1: put port-address into DX register
 - Step 2: put data-value into AL register
 - Step 3: output data-value to port-address

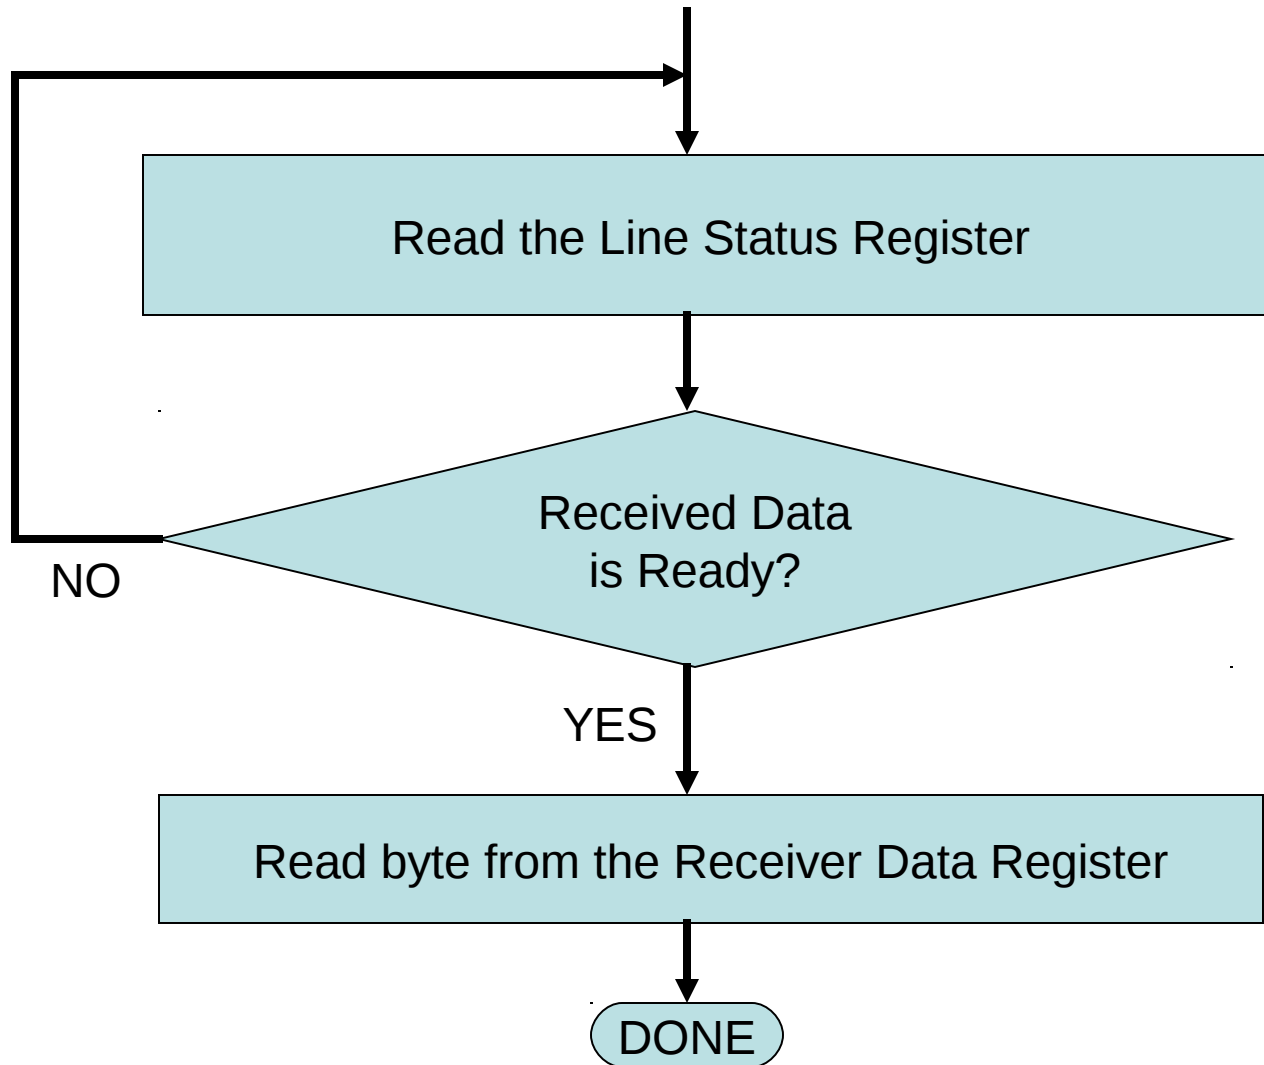
```
mov    $0x03FB, %dx    # UART's Line_Control register
mov    $0x80, %al       # value for setting the DLAB bit
out     %al, %dx        # output data-byte to the i/o-port
```

How 'in' works

- Three steps to perform 'data = in(port);'
 - Step 1: put port-address into DX register
 - Step 2: input from port-address to AL register
 - Step 3: assign value in AL to location 'data'

```
mov    $0x03FD, %dx    # UART's Line_Status register
in      %dx, %al        # input from i/o-port to accumulator
mov     %al, data       # copy accumulator-value to variable
```


How to receive a byte



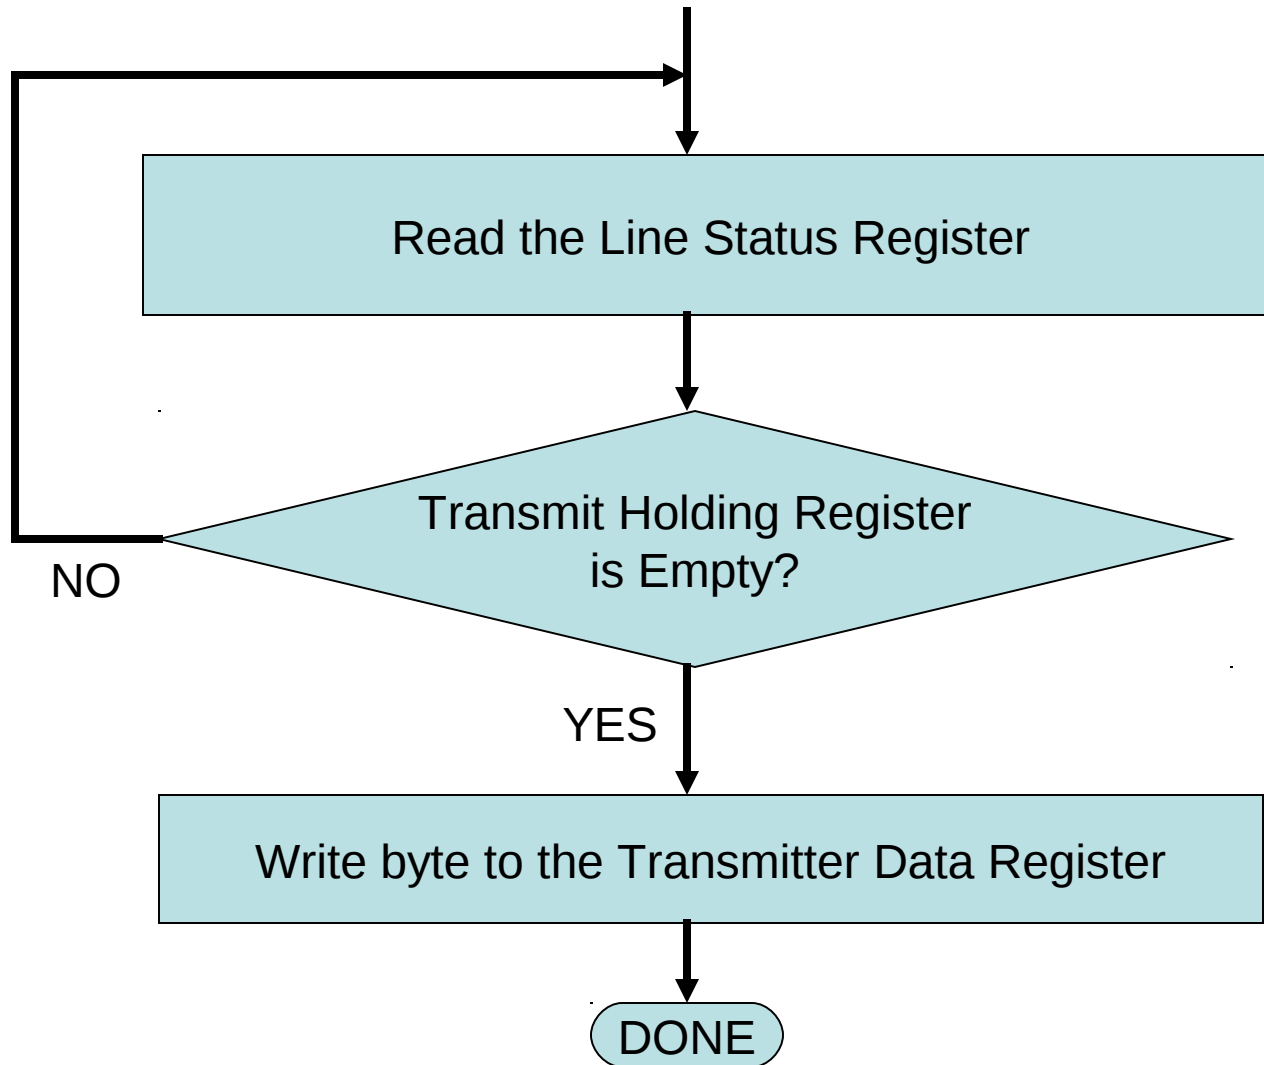
Rx implementation

This assembly language code-fragment inputs a byte of data from
the remote PC, then assigns it to a memory-location labeled 'inch'

```
spin1:  mov    $0x03FD, %dx    # UART's Line_Status register
        in     %dx, %al       # input the current line-status
        test   $0x01, %al     # is the RDR-bit zero?
        jz     spin1         # yes, no new data received

        mov    $0x03F8, %dx    # UART's RxData register
        in     %dx, %al       # input the new data-byte
        mov    %al, inch      # store the data into 'inch'
```

How to transmit a byte



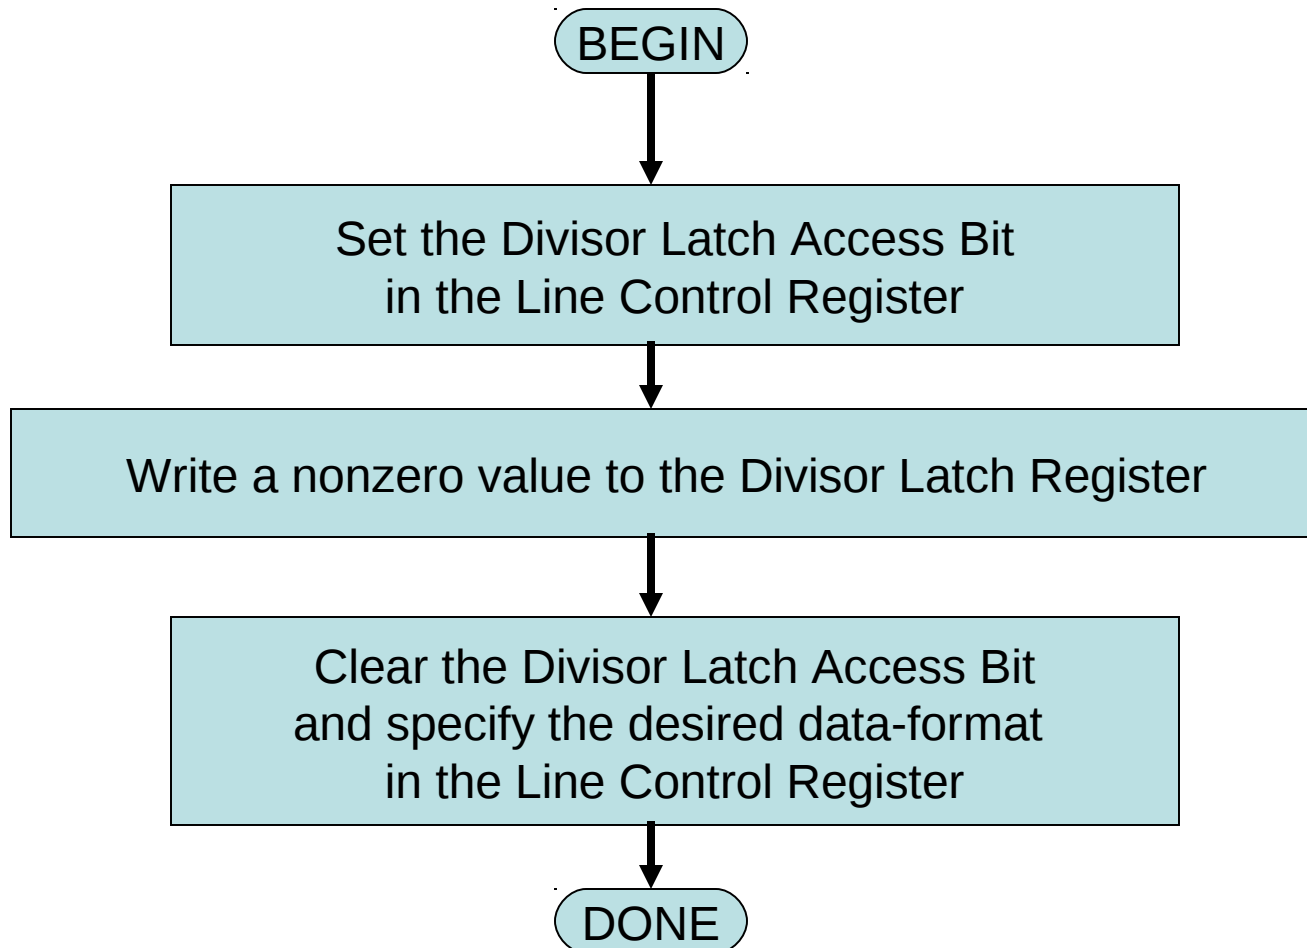
Tx implementation

This assembly language code-fragment fetches a byte of data from
a memory-location labeled 'outch', then outputs it to the remote PC

```
spin2:    mov     $0x03FD, %dx    # UART's Line_Status register
          in      %dx, %al        # input the current line-status
          test    $0x20, %al      # is the THRE-bit zero?
          jz      spin2           # yes, transmitter is still busy

          mov     $0x03F8, %dx    # UART's TxData register
          mov     outch, %al      # get the value to transmit
          out     %al, %dx        # output byte to TxD register
```

Initializing the UART



Init-UART implementation

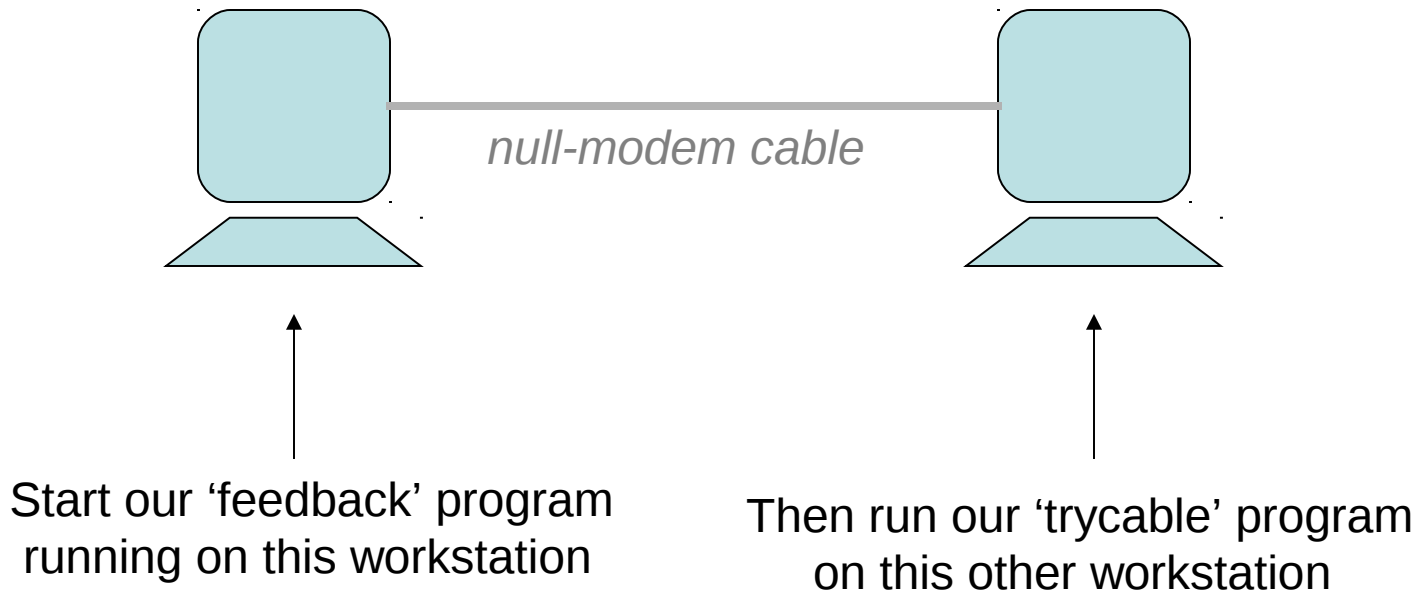
This assembly language code-fragment initializes the UART for
'polled-mode' operation at 115200 baud and 8-N-1 data-format

| | | |
|-----|---------------|---------------------------------|
| mov | \$0x03FB, %dx | # UART Line_Control register |
| mov | \$0x80, %al | # set the DLAB-bit (bit 7) to 1 |
| out | %al, %dx | # for access to Divisor_Latch |
| | | |
| mov | \$0x03F8, %dx | # UART Divisor_Latch register |
| mov | \$0x0001, %ax | # use 1 as the divisor-value |
| out | %ax, %dx | # output the 16-bit latch-value |
| | | |
| mov | \$0x03FB, %dx | # UART Line_Control register |
| mov | \$0x03, %al | # set data-format to 8-N-1 |
| out | %al, %dx | # establish UART data-format |

In-class exercises

- You can try assembling, linking, and then running our 'feedback.s' demo-program (use the 'trycable' program to send some characters via the 'null-modem' cable)
- Can you modify this 'feedback.s' program so that it will continuously display and transmit every character it receives?

Exercise illustration



Now watch what happened on these two screens