

# What is 'bootstrapping'?

The process of loading larger programs for execution than can fit within a single disk-sector

# Serious Pentium explorations

- Experimenting with most Pentium features will require us to write larger-size demos than can fit in one disk-sector (512 bytes)
- So we need a way to load such programs into memory when no OS is yet running
- And we need a convenient way to place such programs onto a persistent storage medium so they can easily be accessed

# Our classroom setup


- Our workstations' hard disks have been 'partitioned' in way that provides a large unused storage-area for us to use freely
- But other portions of these hard disks are dedicated to supporting vital courseware for students who are taking other classes
- We have to understand how to access our 'free' area without disrupting anyone else

# Fixed-Size 'blocks'

- All data-transfers to and from the hard disk are comprised of fixed-size blocks called 'sectors' (whose size equals 512 bytes)
- On modern hard disks, these sectors are identified by sector-numbers starting at 0
- This scheme for addressing disk sectors is known as Logical Block Addressing (LBA)
- So the hard disk is just an array of sectors

# Visualizing the hard disk

**A large array of 512-byte disk sectors**



**Disk storage-capacity (in bytes) = (total number of sectors) x (512 bytes/sector)**

Example: If disk-capacity is 160 GigaBytes, then the total number of disk-sectors can be found by division:

$(160000000000 \text{ bytes}) / (512 \text{ bytes-per-sector})$

assuming that you have a pocket-calculator capable of displaying enough digits!


# Disk Partitions

- The total storage-area of the hard disk is usually subdivided into non-overlapping regions called 'disk partitions'



# Master Boot Record


- A small area at the beginning of the disk is dedicated to 'managing' the disk partitions



- In particular, sector number 0 is known as the Master Boot Record (very important!)

# Format of the MBR

- The MBR is subdivided into three areas:
  - The boot loader program (e.g., GRUB)
  - The 'partition table' data-structure
  - The MBR signature (i.e., 0x55, 0xAA)



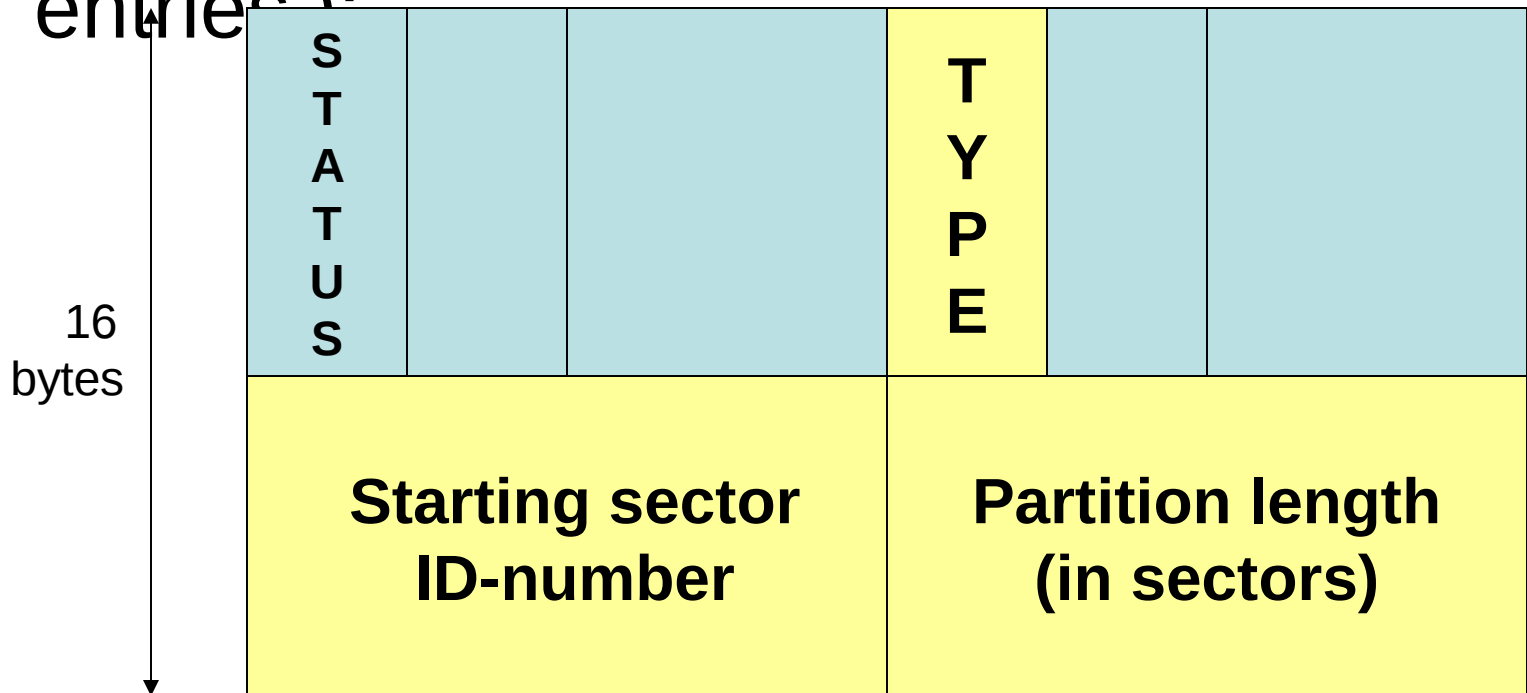


# ‘Reading’ the MBR

- To see the hard disk’s Partition Table, we must ‘read’ the entire Master Boot Record
- (We ignore the boot-loader and signature)
- But we will need to understand the format of the data stored in that Partition Table
- We first need to know how to devise code that can transfer the MBR (sector 0) from the hard-disk into a suitable memory-area

# Partition Table Entries

- The MBR is an array containing four data-structures (called 'partition-table entries').



Some fields contain 'obsolete' information

# TYPE-ID

- Each partition-table entry has a TYPE-ID
  - TYPE-ID is 0x07 for a 'Windows' partition
  - TYPE-ID is 0x83 for our 'Linux' partition
  - TYPE-ID is 0x00 when the entry is 'unused'
- You can find a list of TYPE-ID numbers posted on the internet (see our website)
- Our disks have an extra 'Linux' partition that nobody else is using this semester

# BIOS Disk Drive Services

- An assortment of disk-access functions is available under software Interrupt 0x13
- Originally there were just six functions (to support IBM-PC floppy diskette systems)
- More functions were added when PC/XTs introduced the use of small Hard Disks
- Now, with huge hard disk capacities, there is a set of “Enhanced Disk Drive” services

# Phoenix Technologies Ltd

- You can find online documentation for the BIOS EDD specification 3.0 (see website)
- We'll use function 0x42 to read the MBR
- It requires initializing some fields in a small data-structure (the "Disk-Address Packet")
- Then we load parameters in four registers (DS:SI = address of the DAP, DL = disk-ID and AH = 0x42) and execute 'int \$0x13'

# EDD Disk-Address Packet

7	6	5	4	3	2	1	0
segment-address of transfer-area		offset-address of transfer area		reserved (=0x00)	sector count	reserved (=0x00)	packet length
Logical Block Address of disk-sector (64-bits)							
Physical-address of memory transfer-area (64-bits) (in case segment:offset above is 0xFFFF:FFFF)							

# The MBR parameters

Here are assembly language statements that you could use to create a Disk Address Packet for reading the hard-disk's Master Boot Record into the memory-area immediately following the 512-byte BOOT\_LOCN area

```
#-----  
packet: .byte    16, 0          # packet-size = 16 bytes  
        .byte    1, 0          # sector-count = 1 sector  
        .word    0x0200, 0x07C0 # transfer-area's address  
        .quad    0             # MBR's Logical Block Address  
#-----
```

Our demo-program (named 'finalpte.s') uses statements similar to these.

# How we search the Partition Table

- Our demo-program 'finalpte.s' locates the final valid entry in the MBR Partition Table
- It displays the contents of that entry (i.e., four longwords) in hexadecimal format
- To do its search, it simply scans the table entries in backward order looking for the first entry that has a nonzero 'type' code



# Instructions we could use

```

    mov     $0x03FE, %si      # point DS:SI to signature-word
    mov     $4, %cx          # setup count of table's entries
nxpte:
    sub     $16, %si          # back up to the previous entry
    cmpb    $0x00, 4(%si)     # entry's type-code is defined?
    loope   nxpte             # no, examine the next entry

    jcxz     nopte            # search fails if CX reached 0

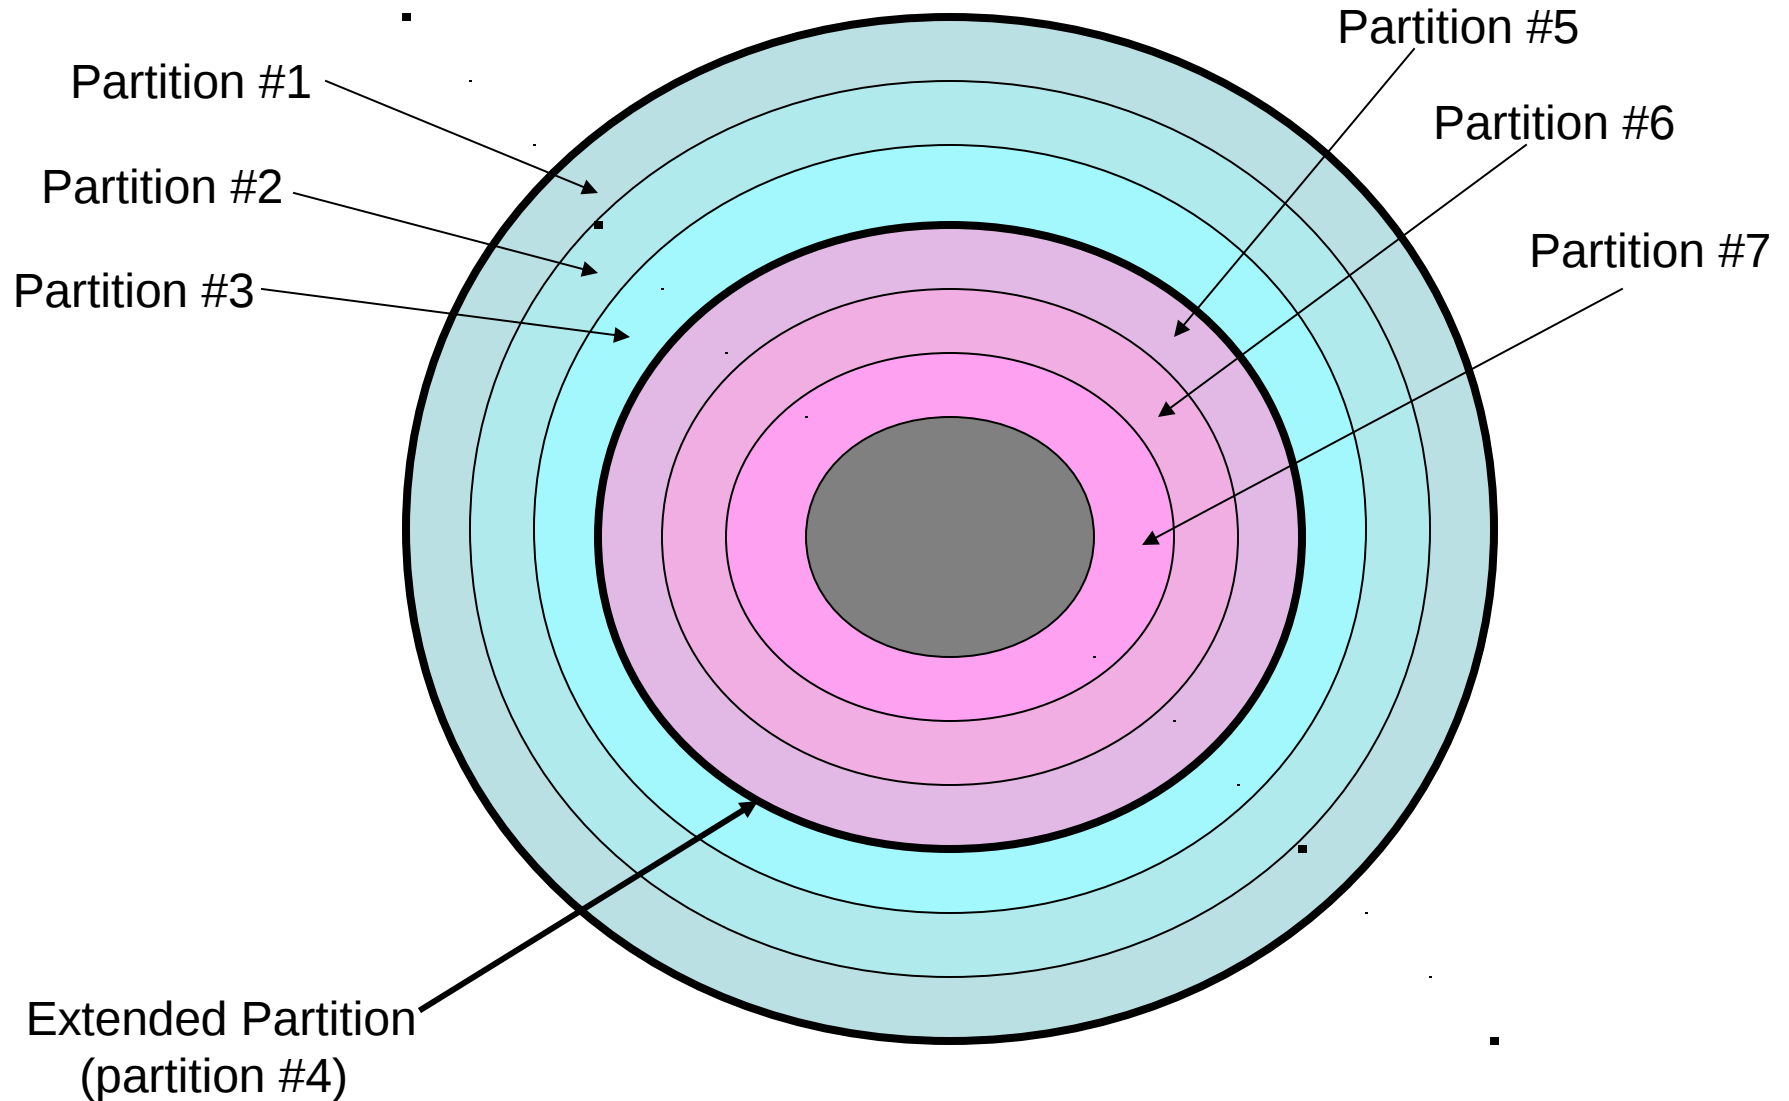
    # If we get here, then DS:SI is pointing to the final valid PT-entry
    jmp     found

nopte:  # We should never arrive here -- unless no valid partitions exist
```

# ‘Extended’ partitions

- The hard-disk’s Master Boot Record only has room for four Partition-Table Entries
- But some systems need to use more than four disk-partitions, so a way was devised to allow one of the MBR’s partition-table entries to describe an ‘extended’ partition
- This scheme does not seem to have been ‘standardized’ yet -- hence, confusion!

# The Linux scheme



# Our 'cs686ipl.s' boot-loader

- We created a 'boot-loader' that will work with Linux systems that have 'extended' disk-partitions ("Initial Program Loader")
- It uses the EDD Read\_Sectors function, and it will read 127 consecutive sectors from the disk's final Linux-type partition
- It transfers these disk-sectors into the memory-arena at address 0x00010000

# Our 'controls.s' demo-program

- To demonstrate our boot-loader, we wrote a short program that can be 'loaded' into memory at 0x10000 and then executed
- It will display some useful information (and thereby verify that our boot-loader worked)
- Our boot-loader requires that a special 'program signature' (i.e., 0xABCD) must occupy the first word of any program it attempts to execute (as a 'sanity' check)

# Depiction of 'boot-strapping'

We install our 'cs686ipl.b' loader into the boot-sector of disk-partition number 4:

```
$ dd if=cs686ipl.b of=/dev/sda4
```

We install our 'controls.b' demo-program into the subsequent disk-sectors:

```
$ dd if=controls.b of=/dev/sda4 seek=1
```

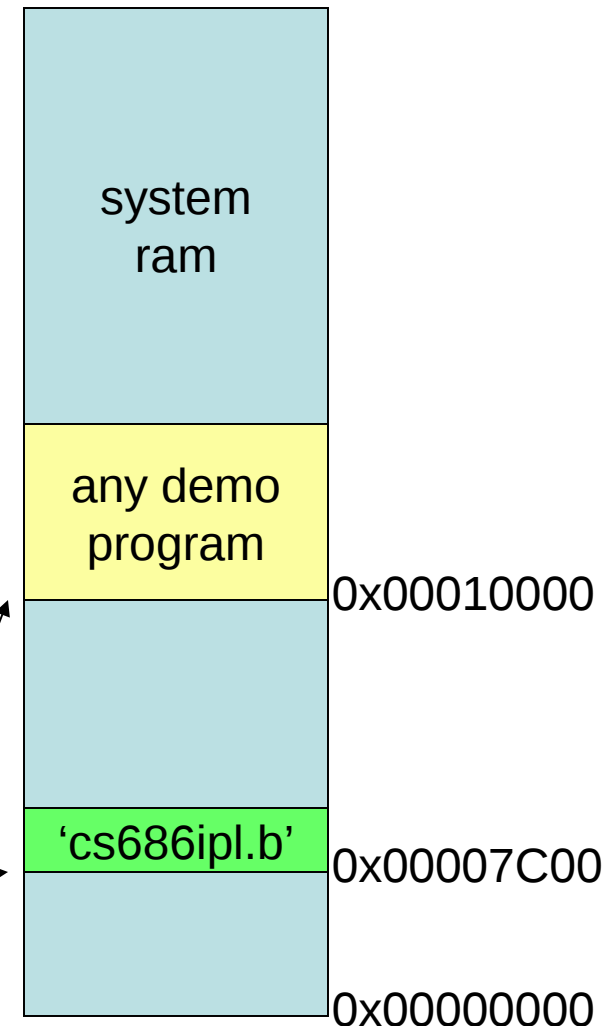
Then we 'reboot' our machine to begin the bootstrapping process...

Step 1: The ROM-BIOS firmware loads GRUB

Step 2: The user selects a disk-partition from the GRUB program's menu's options

Step 3: GRUB loads our 'cs686ipl.b' loader

Step 4: 'cs686ipl.b' loads our program demo



# In-class exercise #1

- Install the 'cs686ipl.b' boot-loader on your assigned 'anchor' machine:

```
$ dd if=cs686ipl.b of=/dev/sda4
```

- Also install the 'controls.b' demo-program on your assigned 'anchor' machine:

```
$ dd if=controls.b of=/dev/sda4 seek=1
```

- Then use the 'fileview' utility-program (from our class website, under 'Resources') to view the first few disk-sectors in partition number 4:

```
$ fileview /dev/sda4
```

# In-class exercise #2

- Try 'rebooting' your 'anchor' machine, to see the information shown by 'controls.b'
  - Use 'ssh' to log on to the 'colby' gateway
  - Use 'telnet' to log onto your 'anchor' machine
  - Use the 'sudo reboot' command to reboot and then watch for the GRUB menu-selection that will launch the 'boot-strapping' process which will execute the demo-program you installed



# In-class exercise #3

- Can you apply what you've learned in our prior exercises to enhance our 'controls.s' demo-program so that it will display some additional information about the CPU's register-values at boot-time?
- For example, could you display the values held in the GDTR and IDTR registers? (Remember those are 48-bit registers)