


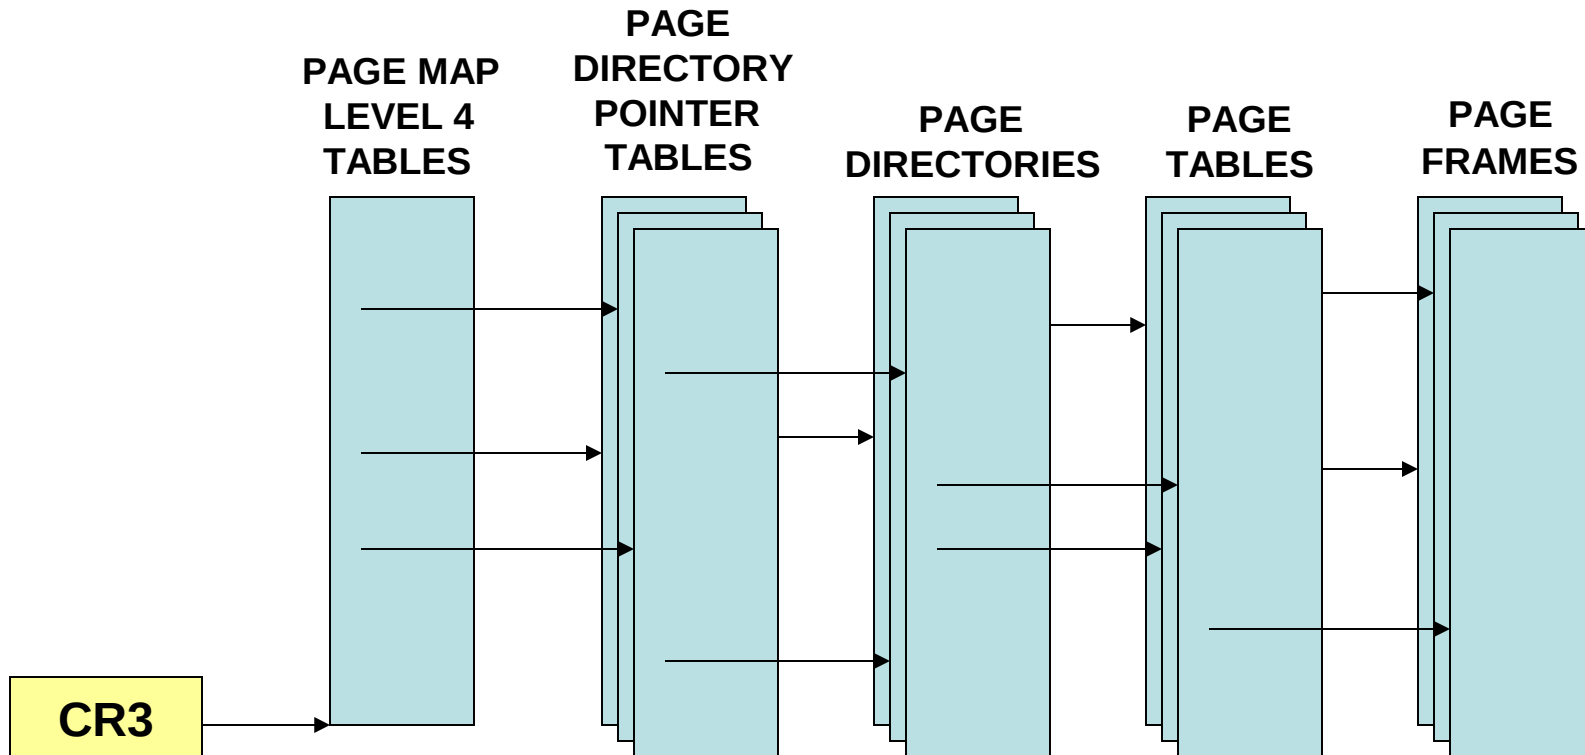
# Venturing into 64-bit mode

Examining the steps needed to  
take the processor into IA-32e  
mode -- and then back out again

# Two-Level Translation Scheme

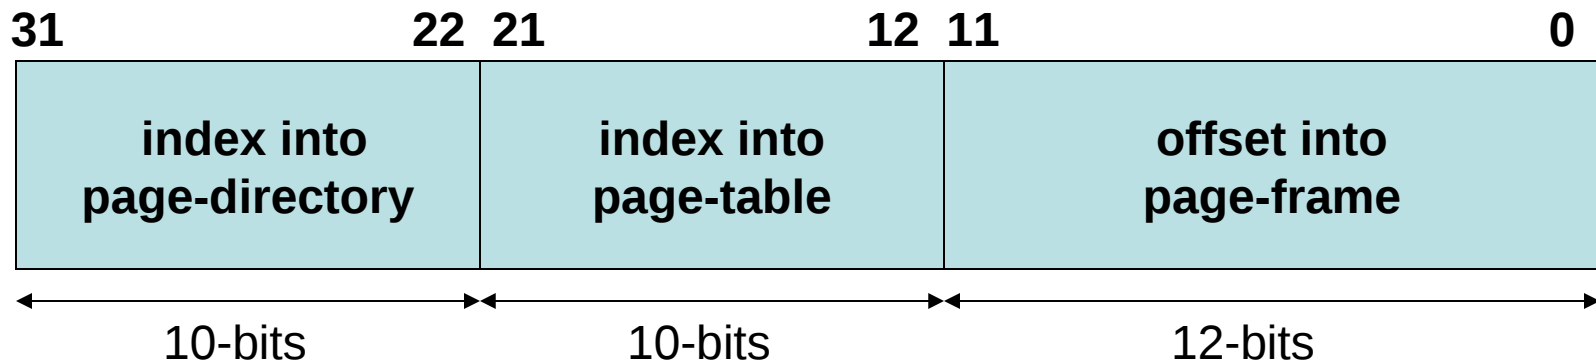


# Four-Level Translation Scheme



# 2-level address-translation

- The CPU examines any virtual address it encounters, subdividing it into three fields



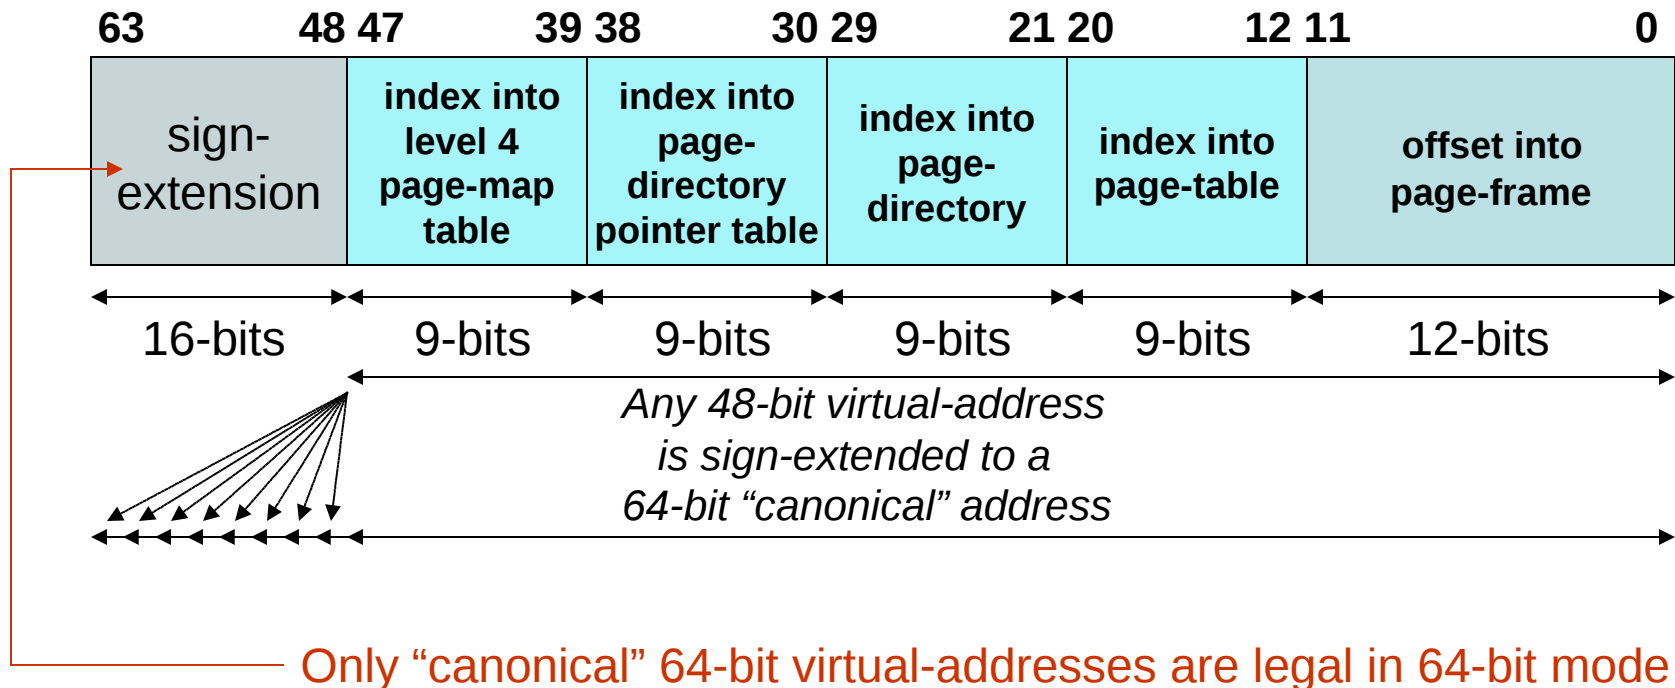
This field selects one of the 1024 array-entries in the Page-Directory

This field selects one of the 1024 array-entries in that Page-Table


This field provides the offset to one of the 4096 bytes in that Page-Frame

# 4-level address-translation

- The CPU examines any virtual address it encounters, subdividing it into **five** fields



# Format of a Page-Directory entry



## LEGEND

P = Present (1=yes, 0=no)

W = Writable (1 = yes, 0 = no)

U = User (1 = yes, 0 = no)



A = Accessed (1 = yes, 0 = no)

PS = Page-Size (0=4KB, 1 = 4MB)

PWT = Page Write-Through (1=yes, 0 = no)

PCD = Page Cache-Disable (1 = yes, 0 = no)

# Format of 64-bit table-entries



Meaning of —————  
these bits varies with the table

## Legend:

P = Present (1=yes, 0=no)  
 W = Writable (1=yes, 0=no)  
 U = User-page (1=yes, 0=no)  
 A = Accessed (1=yes, 0=no)

PWT = Page Cache Disable (1=yes, 0=no)  
 PWT = Page Write-Through (1=yes, 0=no)  
 avl = available for user-defined purposes  
 EXB = Execution-disabled Bit (if EFER.NXE=1)

# Extended Feature Enable Register

- This Model-Specific Register (MSR) was introduced in the AMD64 architecture and perpetuated by EM64T (for compatibility)



## Legend:

- SCE = SysCall/sysret is Enabled (1=yes, 0=no)
- LME = Long-Mode is Enabled (1=yes, 0=no)
- LMA = Long-Mode is Active (1=yes, 0=no)
- NXE = Non-eXecutable pages Enabled (1=yes, 0=no)

**NOTE:** The MSR address-index for EFER = 0xC0000080, and this register is accessed using RDMSR or WRMSR instructions



# RDMSR and WRMSR


- An assembly language code-fragment to turn on the LME-bit ('Long-Mode' Enable):

```
# Each Model-Specific Register (MSR) is 64-bits wide and has a unique  
# 32-bit address-index which is first placed into register ECX. Then the  
# least-significant 32-bits of that MSR is accessed using register EAX,  
# while the most-significant 32-bits is accessed using register EDX.
```


```
mov      $0xC0000080, %ecx      # setup EFER address-index  
rdmsr                                # read EFER into (EDX,EAX)  
bts      $8, %eax              # set the LME-bit's image to 1  
wrmsr                                # write (EDX,EAX) into EFER
```

```
# NOTE: RDMSR and WRMSR must be executed at 'Ring0' privilege-level.
```

# Control Registers CR4 and CR0



Control Register CR4



Control Register CR0

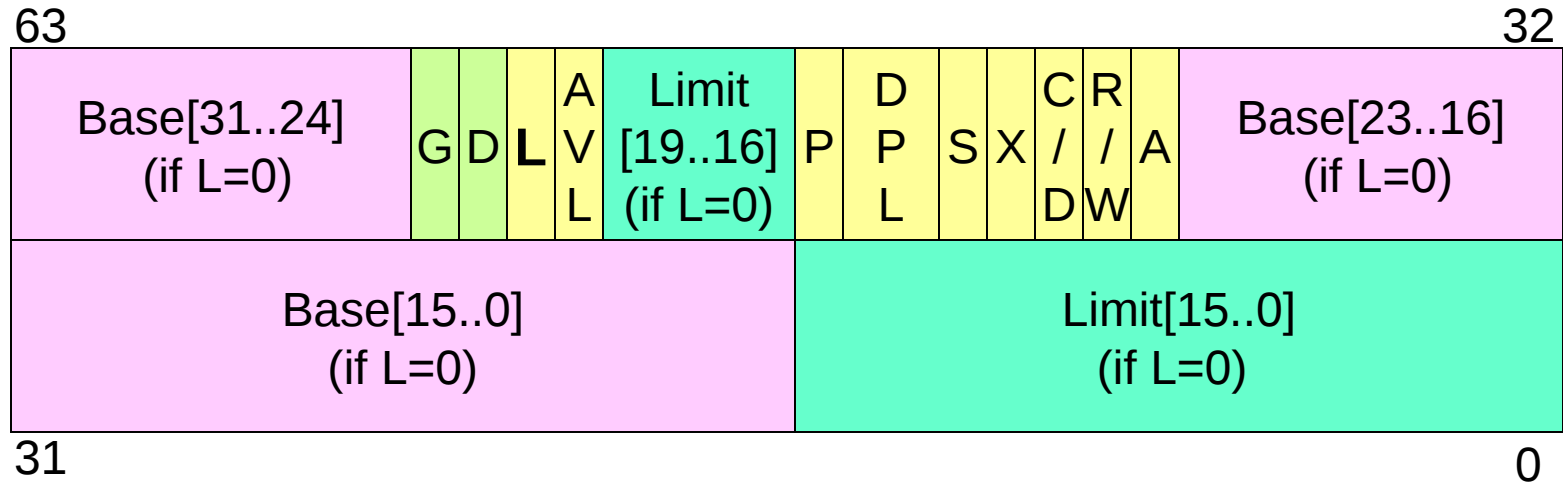
## Legend (for 64-bit mode):

PE = Protected-mode Enabled (1=yes, 0=no)

PG = Paging Enabled (1=yes, 0=no)

PAE = Page-Addressing Extensions (1=enabled, 0=disabled)

# Segment-Descriptor Format



## Legend:

G = Granularity (0 = byte, 1 = 4KB-page)

D = Default size (0 = 16-bit, 1 = 32-bit)

X = eXecutable (0 = no, 1 = yes)

DPL = Descriptor Privilege Level (0..3)

P = Present (0 = no, 1 = yes)

S = System (0 = yes, 1 = no)

A = Accessed (0 = no, 1 = yes)

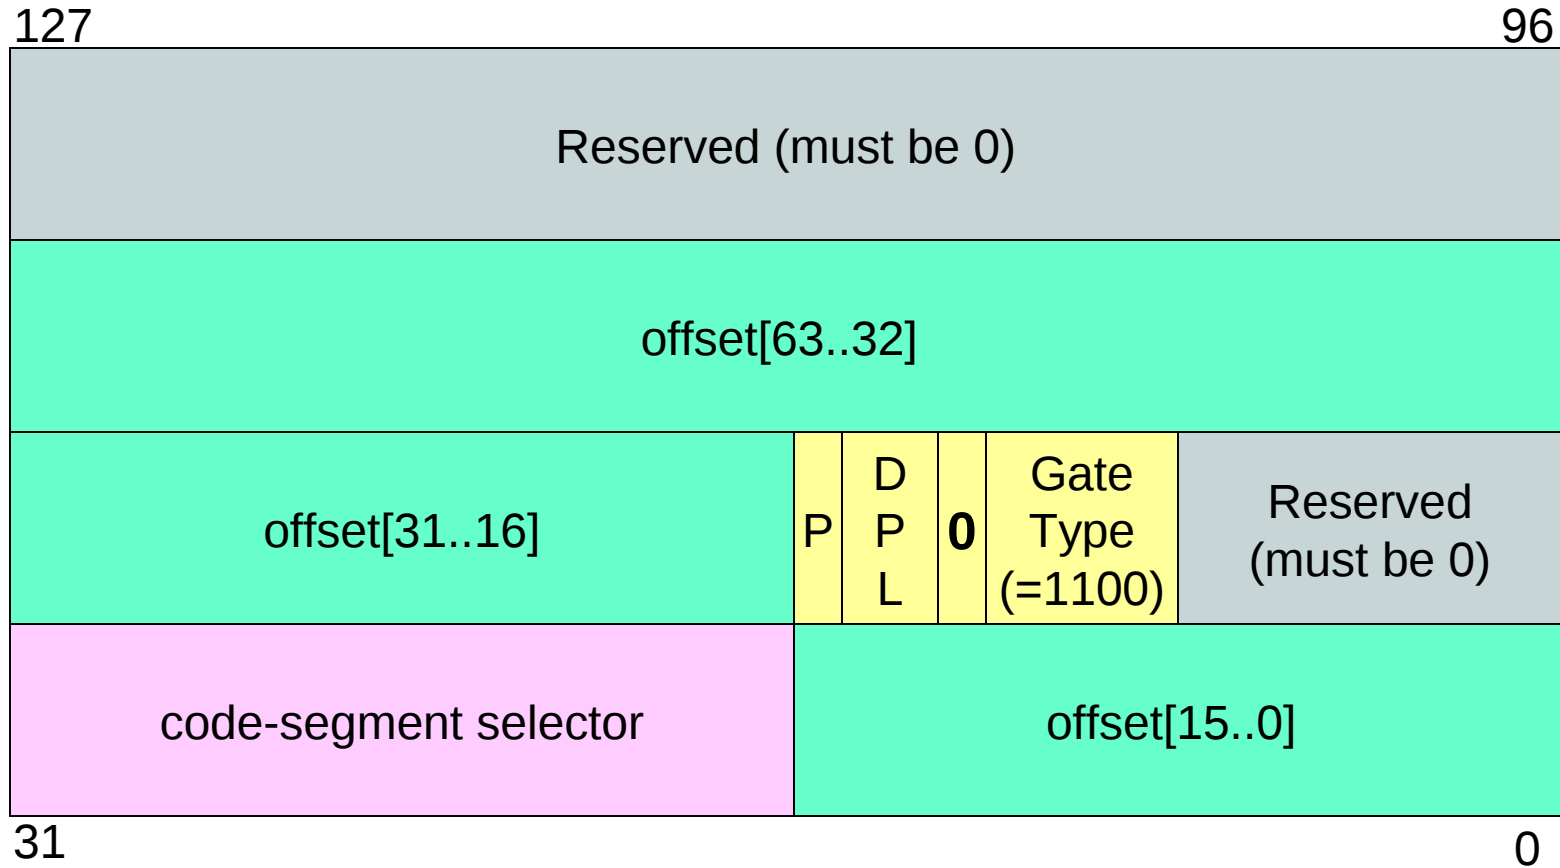
*code-segments:* R = Readable (0 = no, 1 = yes) C = Conforming (0=no, 1=yes)

*data-segments:* W = Writable (0 = no, 1 = yes) D = expands-Down (0=no, 1=yes)

L = Long-mode (i.e., 64-bit addressing) (0=no, 1=yes)

AVL = Available for user's purposes

# IA-32e Call-Gate descriptor



*We can use a call-gate to 'jump' from 16-bit code-segment to a 64-bit code-segment*

# Summary of steps

- Transition from real-mode to IA-32e mode:
  - Build the table of global descriptors
  - Load GDTR with pseudo-descriptor for GDT
  - Build the 4-level page-mapping tables
  - Enable IA-32e mode (set EFER.LME=1)
  - Enable Page-Address Extensions (CR4.PAE)
  - Load Level4 page-map table address in CR3
  - Activate IA-32e mode (CR0.PE and CR0.PG)
  - Transfer via call-gate to 64-bit code-segment

# Notes on the transition

- Code-segment must be “identity-mapped”
- Interrupts have to be temporarily disabled
- All memory-addressing in 64-bit mode via CS, SS, DS or ES uses 0 as base-address (and checking of segment-limits is omitted)

# For a return to 'real-mode'

- Processor must enter 16-bit code-segment in 'compatibility-mode' via indirect far jump
  - Load segment-registers DS, ES, and SS with 'writable' 16-bit segment-selectors (64K-limit)
  - Code-segment has to be "identity-mapped"
  - Deactivate IA-32e mode by clearing PG-bit
  - Leave 'protected-mode' by clearing PE-bit
  - Reload registers CS and SS with real-mode segment-addresses before enabling interrupts


# Demo-program: 'inandout.s'

- We created a demo-program that starts in 'real-mode', enters 64-bit mode and draws a message, jumps to 'compatibility mode' and draws another message, then returns to real-mode and shows a final message
- It has to write directly to VRAM when it's not executing in real-mode – because the ROM-BIOS routines use real-style code





# How text-mode VRAM works

- The video memory resides at 0x000B8000 and in text-mode it is organized as a linear array of two-byte elements (i.e., ‘words’):




- Array-elements are arranged in “row-major” order (left-to-right, top-to-bottom)

# Default color-programming



# Character-cell screen-locations



# In-class exercise #1

- Can you modify the message-colors used in our 'inandout.s' demo-program so that:
  - the first message is bright-red against white
  - the second message is brown against cyan
  - The final message is magenta against black

# In-class exercise #2

- Can you modify the 'inandout.s' program so that it 'skips' entering 64-bit mode: i.e., so it transfers directly from real-mode to 'compatibility' mode, prints its confirmation message, and then transfers back to 'real-mode' to show a final message?

