


# Interrupts in Protected-Mode

Writing a protected-mode  
interrupt-service routine for the  
timer-tick interrupt

# Rationale

- Usefulness of a general-purpose computer is dependent on its ability to interact with various peripheral devices attached to it (e.g., keyboard, display, disk-drives, etc.)
- Devices require a prompt response from the cpu when various events occur, even when the cpu is busy running a program
- The x86 interrupt-mechanism provides this

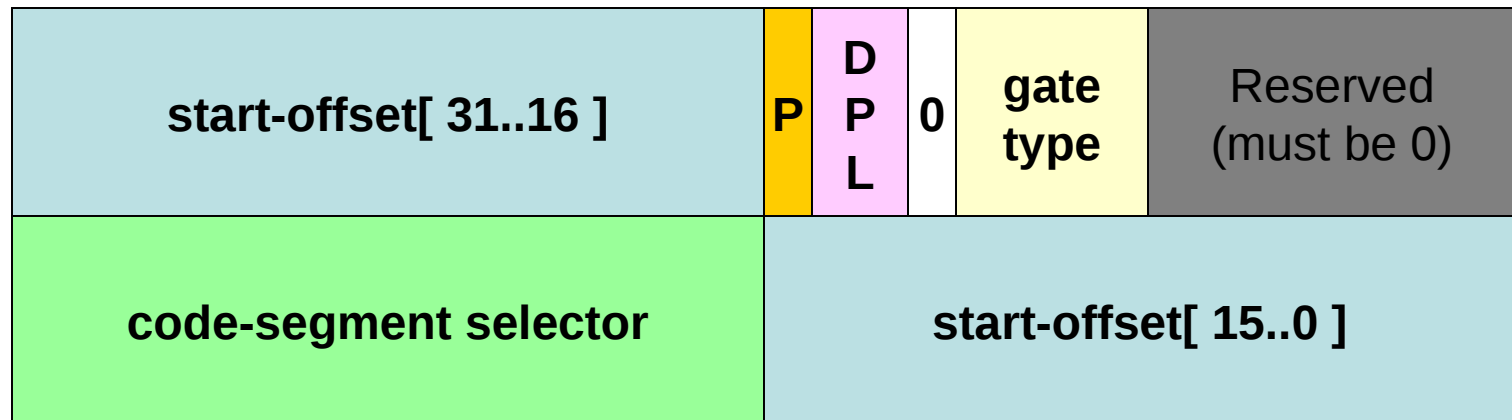
# Simplified Block Diagram



# PM requirements

- Unlike real-mode, where all code executes with full privileges (i.e., ring 0), protected-mode code usually is executed with some privilege restrictions (e.g., usually ring3)
- Normally these restrictions prevent direct control of any of the peripheral devices
- Thus, when responding to an interrupt in protected-mode, a ring-transition, and an accompanying stack-switch, are involved

# IA-32 Interrupt-Gate Descriptors



Legend:

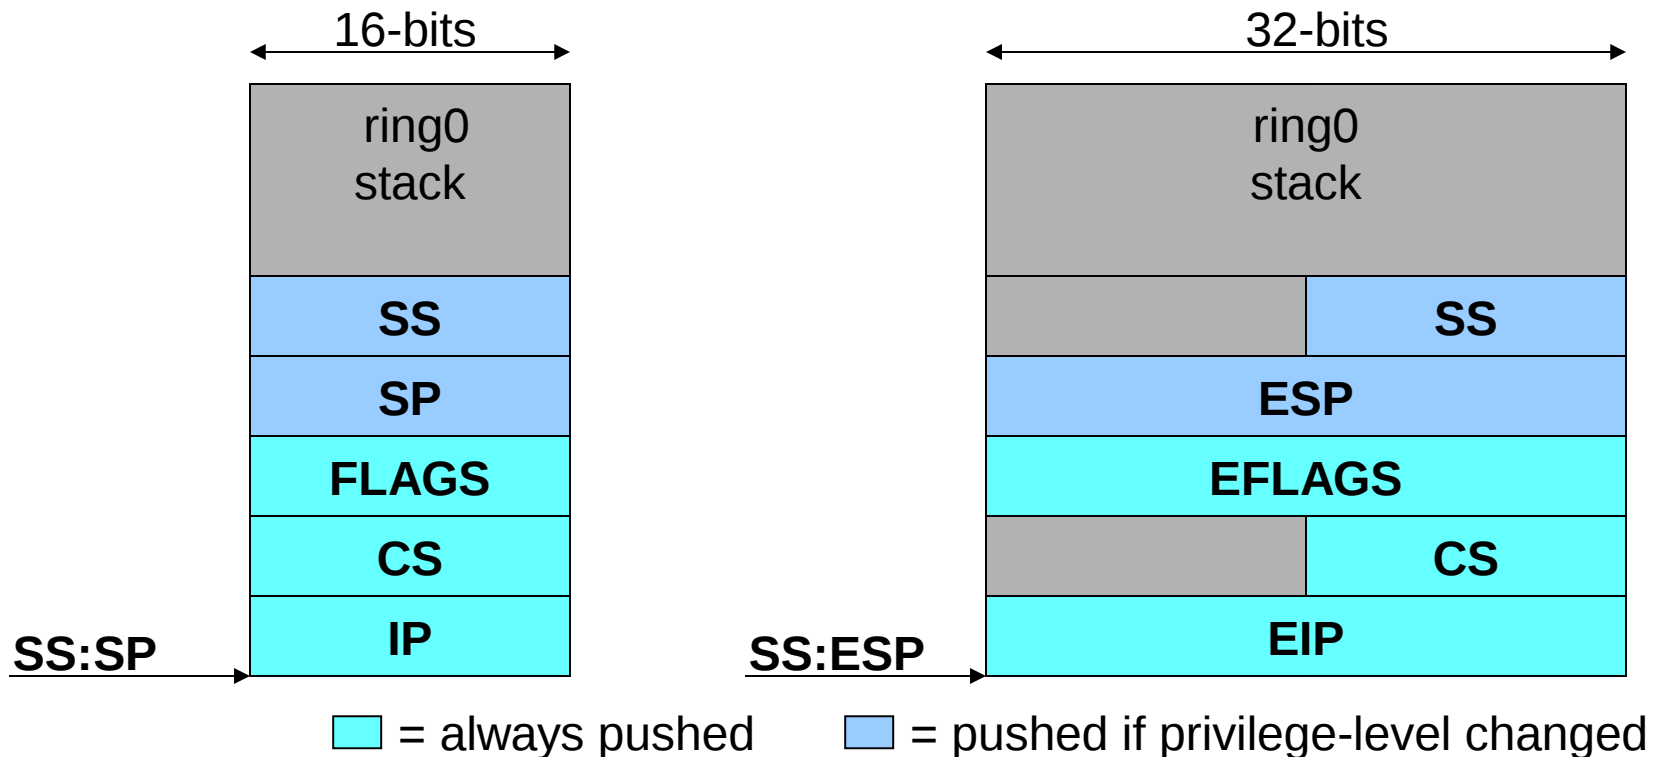
P=present (1=yes, 0=no)      DPL=Descriptor Privilege-Level (0,1,2,3)  
code-selector (specifies memory-segment containing procedure code)  
start-offset (specifies the procedure's entry-point within its code-segment)  
gate-types: 0x6 = 16bit Interrupt-Gate, 0x7 = 16-bit Trap-Gate  
              0xE = 32bit Interrupt-Gate, 0xF = 32-bit Trap-Gate

# Trap-Gate vs. Interrupt-Gate

- The only distinction between a Trap-Gate and an Interrupt-Gate is in whether or not the CPU will automatically clear the IF-bit (Interrupt-Flag in EFLAGS register) as part of its response to an interrupt-request
- This is needed in cases where an Interrupt Service Routine executes outside ring0, so could not execute 'cli' or 'sti' instructions

# 16bit-Gate vs. 32bit-Gate

- The CPU constructs different stackframes for the 16-bit versus the 32-bit gate-types



# Return-from-Interrupt

- The programmer who writes an Interrupt Service Routine must know whether the Gate was 16-bit or 32-bit, in order to use the correct 'interrupt-return' instruction
- In a code-segment whose default-bit is 0 (i.e., .code16), an 'iret' instruction performs the correct return-actions for a 16-bit Gate
- Use 'iretl' for returning with a 32-bit Gate




# Interrupt Descriptor Table

- The Gate-Descriptors for device interrupts form an array (called the IDT) and reside in a special system memory-segment
- The CPU will locate the IDT by referring to the value in its IDTR register (48-bits)
- A pair of special instructions exists which allow reading and writing this register:

`sidt mem ; store IDTR into a memory-operand`

`lidt mem ; load IDTR from a memory-operand`

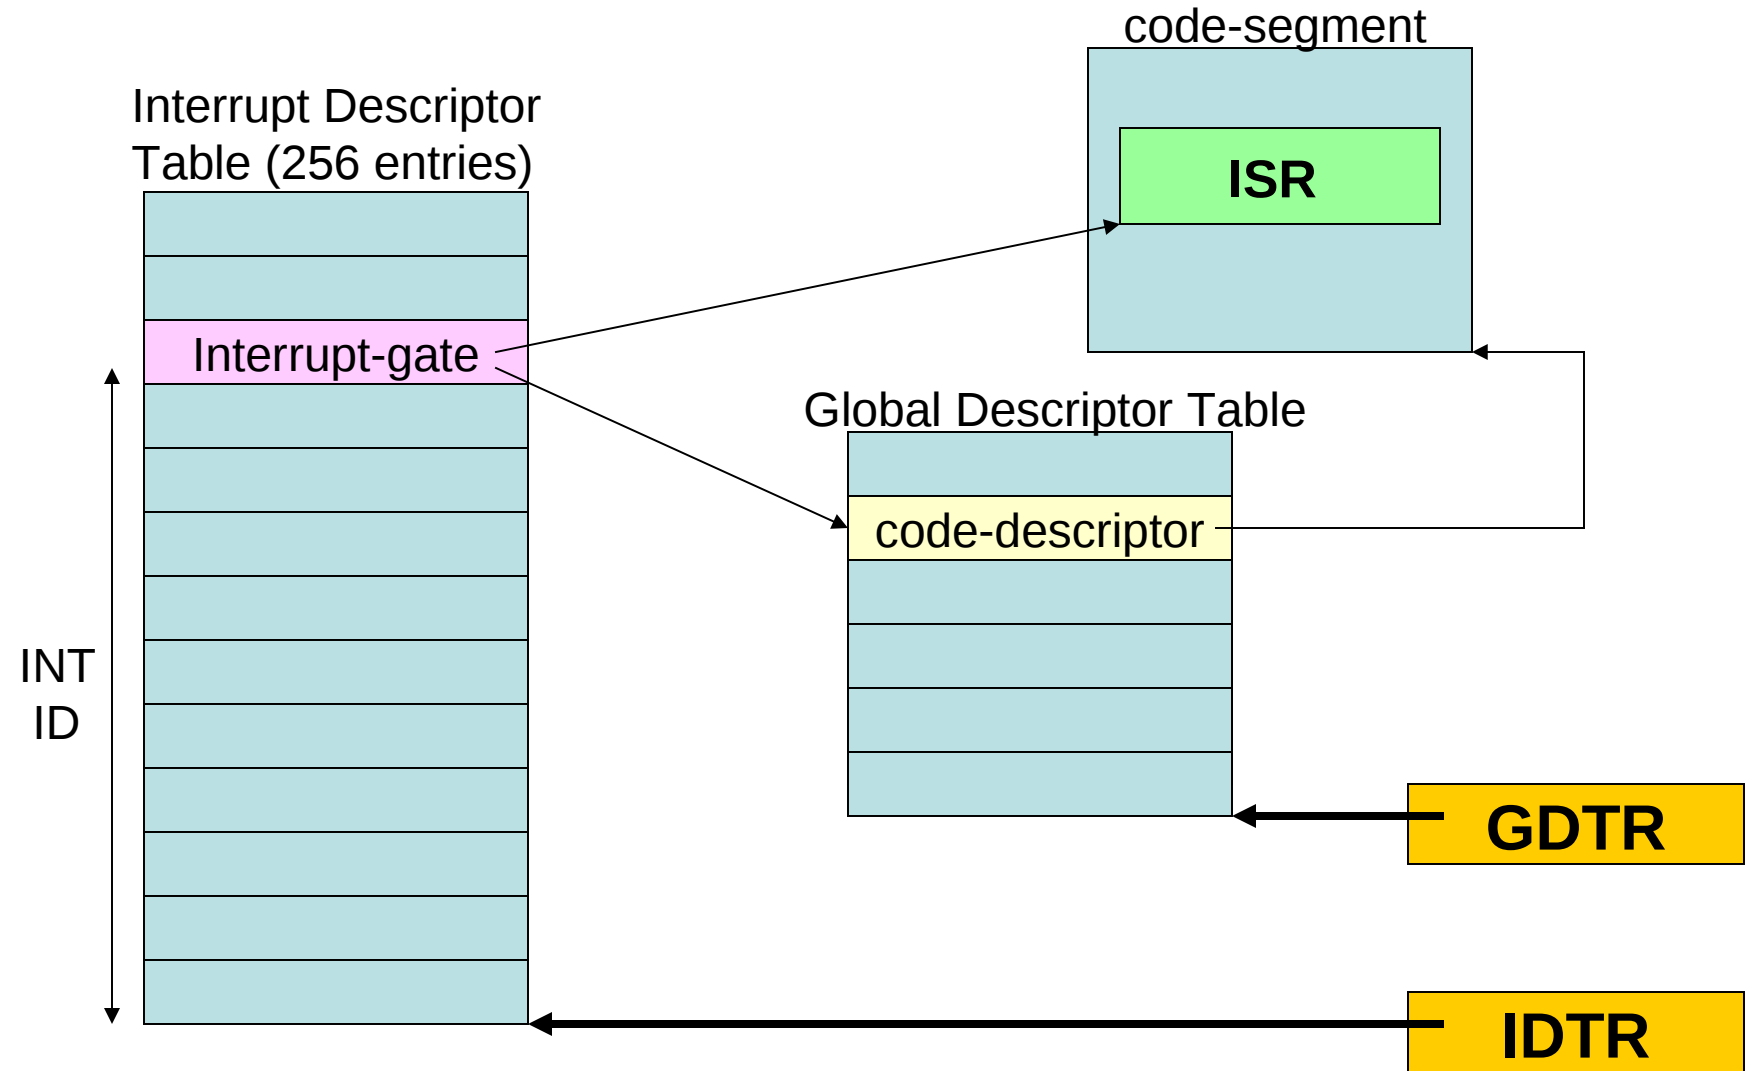
# Format of register IDTR




The instruction '**lidt**' is privileged (can only be executed in ring 0), but the instruction '**sidt**' is unprivileged (it can execute in any ring)

These features are analogous to the instructions '**sgdt**' and '**lgdt**' used to store or to load GDTR (Global Descriptor Table Register)

# Register relationships



# Two Interrupt-Controllers



# Each PIC has a Mask Register

Master PIC  
Interrupt-mask  
(I/O-port 0x21)

|          |          |          |          |          |          |          |          |
|----------|----------|----------|----------|----------|----------|----------|----------|
| IRQ<br>7 | IRQ<br>6 | IRQ<br>5 | IRQ<br>4 | IRQ<br>3 | IRQ<br>2 | IRQ<br>1 | IRQ<br>0 |
|----------|----------|----------|----------|----------|----------|----------|----------|

Slave PIC  
Interrupt-mask  
(I/O-port 0xA1)

|           |           |           |           |           |           |          |          |
|-----------|-----------|-----------|-----------|-----------|-----------|----------|----------|
| IRQ<br>15 | IRQ<br>14 | IRQ<br>13 | IRQ<br>12 | IRQ<br>11 | IRQ<br>10 | IRQ<br>9 | IRQ<br>8 |
|-----------|-----------|-----------|-----------|-----------|-----------|----------|----------|

If a mask-bit is 1, the corresponding device-interrupts are masked;  
If a mask-bit is 0, the corresponding device-interrupts are unmasked

# Demo-program: 'tryisr32.s'

- Let's create a 'protected-mode' program that will handle the timer-tick interrupts
- Its ISR (Interrupt Service Routine) is very similar to the real-mode interrupt-handler
- It increments a 32-bit counter (i.e., 'jiffies')
- It sends EOI-notification to Master PIC
- It must take care to preserve the values that are contained in the CPU's registers

# Two 'threads' in our demo

'main' program-thread


Build descriptor-tables  
Enter protected-mode  
(most interrupts masked)

Do for fifteen seconds:

```
{  
  Convert 'jiffies' value  
    to hours, mins, secs  
  Write time to screen  
}
```

Leave protected-mode  
(most interrupts unmasked)  
Exit to our 'loader'

program-variables (threads share access)



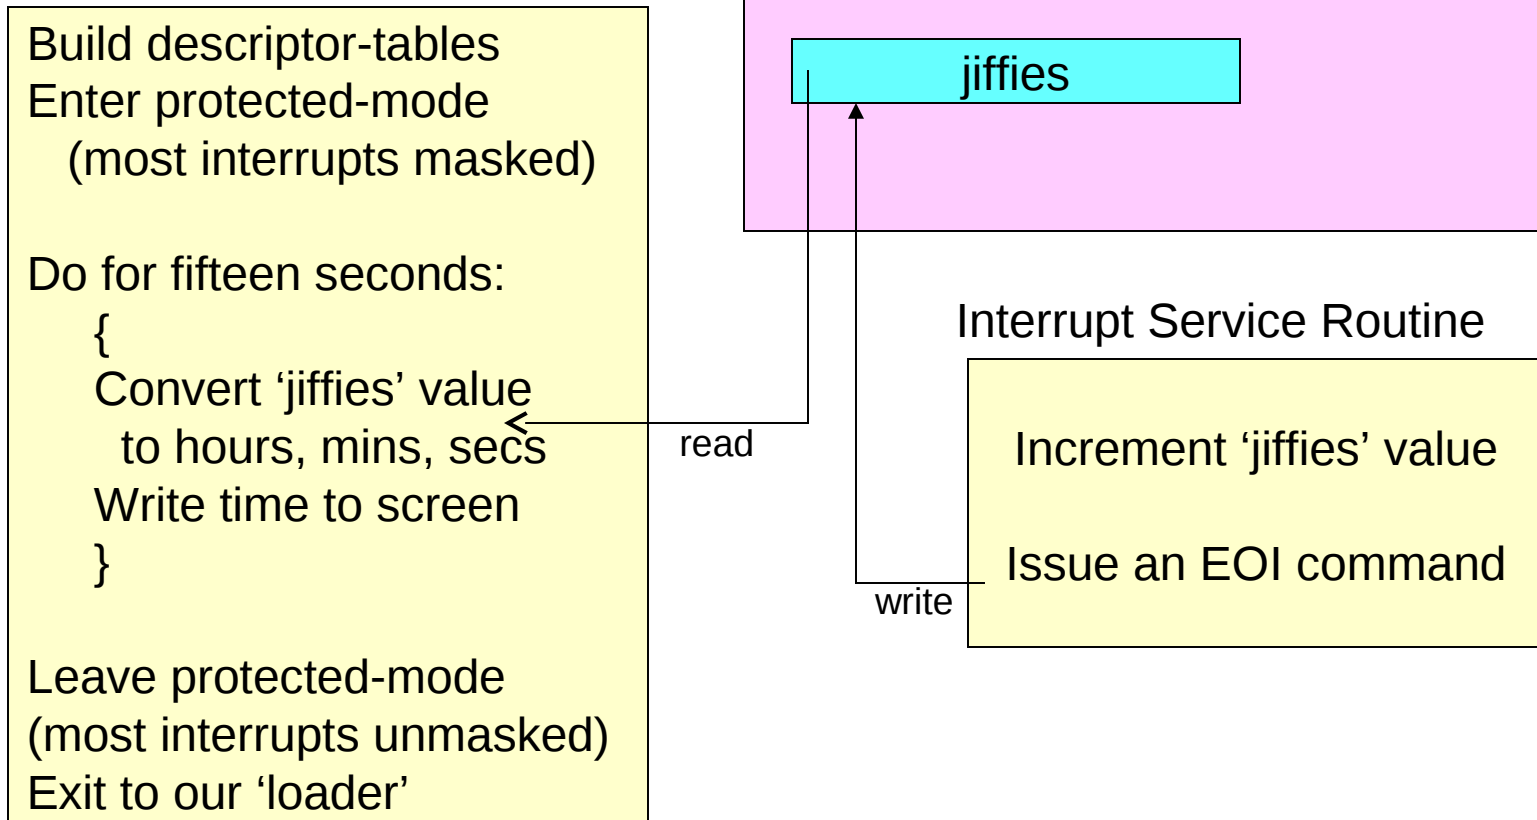
Interrupt Service Routine

Increment 'jiffies' value

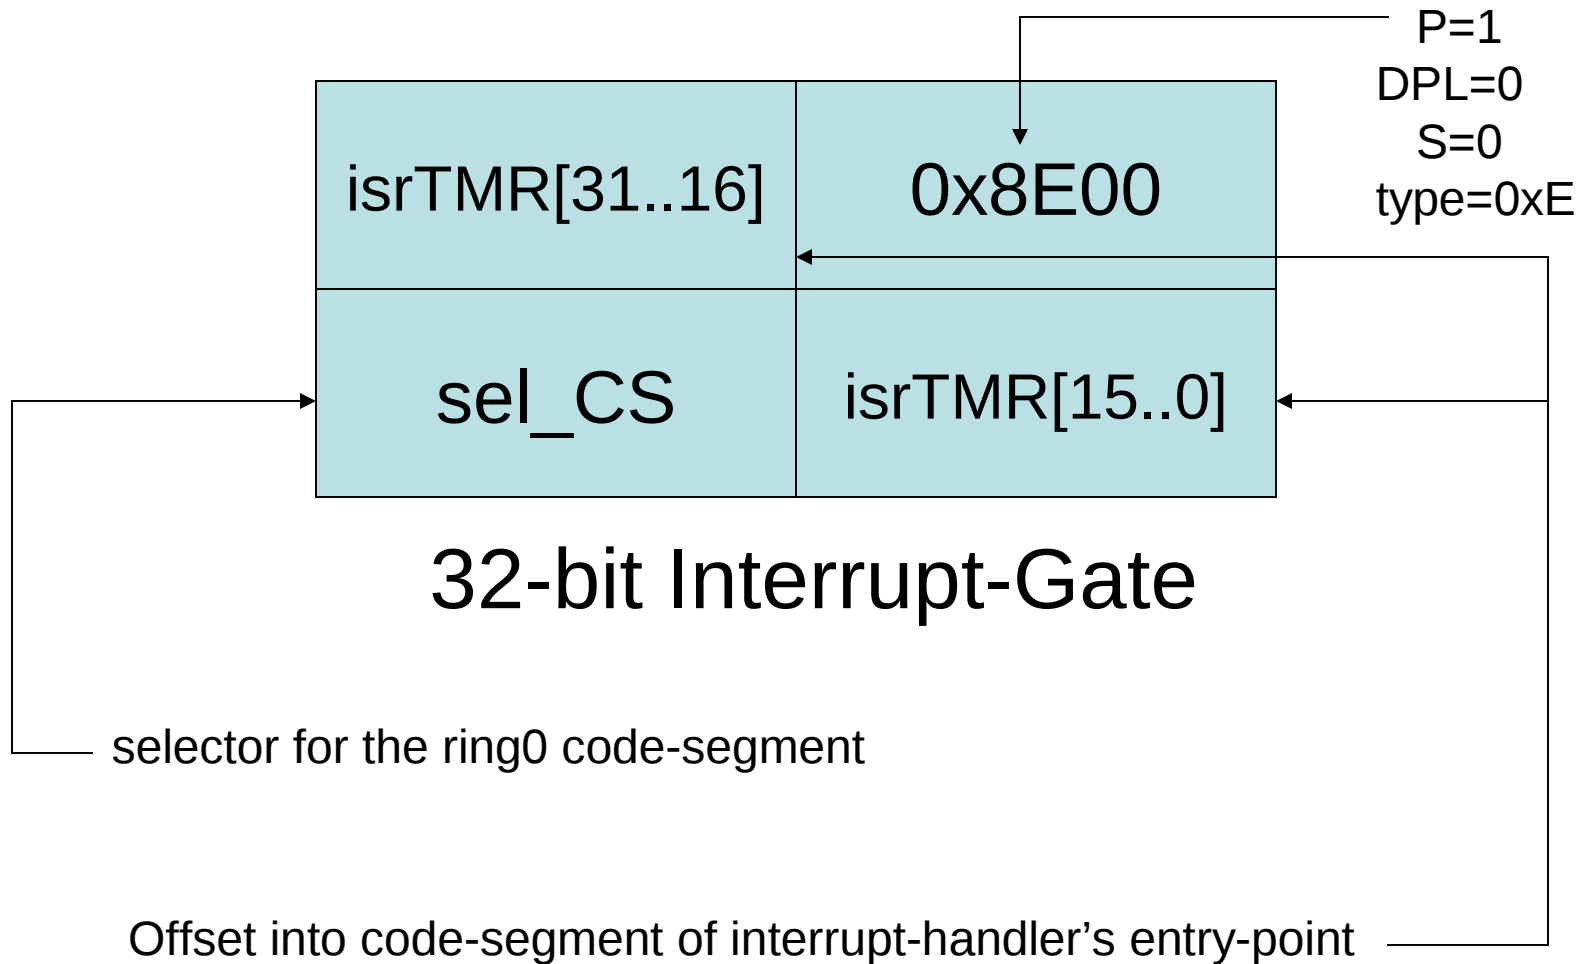
Issue an EOI command

read

write



# Defining INT-8 Gate-Descriptor






# Some Key Steps in the Demo

- Initialize the Descriptor-Table(s)
- Enter Protected-Mode with IF=0
- Load GDTR, IDTR and segment-selectors
- Mask all device-interrupts except timer
- Set IF=1 to enable unmasked interrupts
- Continuously show tick-count (for 15secs)
- Reset IF=0 to disable interrupts (for exit)

# ‘identity-mapping’

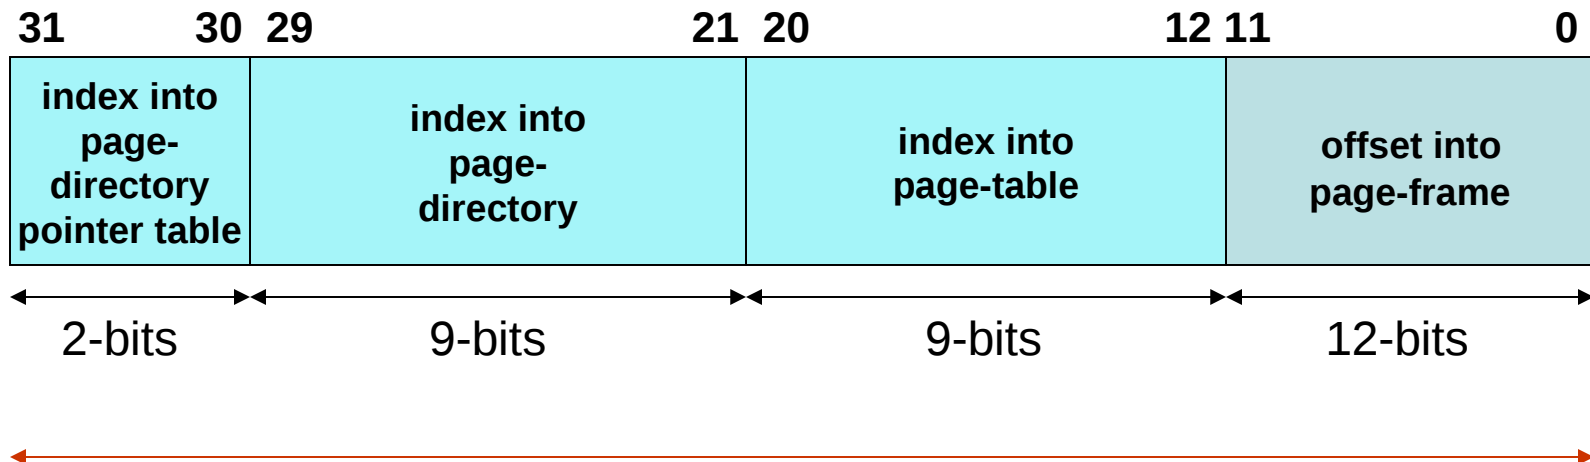
- To make our demo as similar as possible to the code we would need in 64-bit mode (where ‘paging’ is always mandatory), we have chosen to enable paging in our demo and to use Intel’s 3-level mapping scheme (which supports 36-bit physical addresses) because it’s most like the 64-bit mode
- However, we use an ‘identity-mapping’ of the bottom 1-megabyte (to keep it simple)

# 3-Level Translation Scheme



# 3-level address-translation

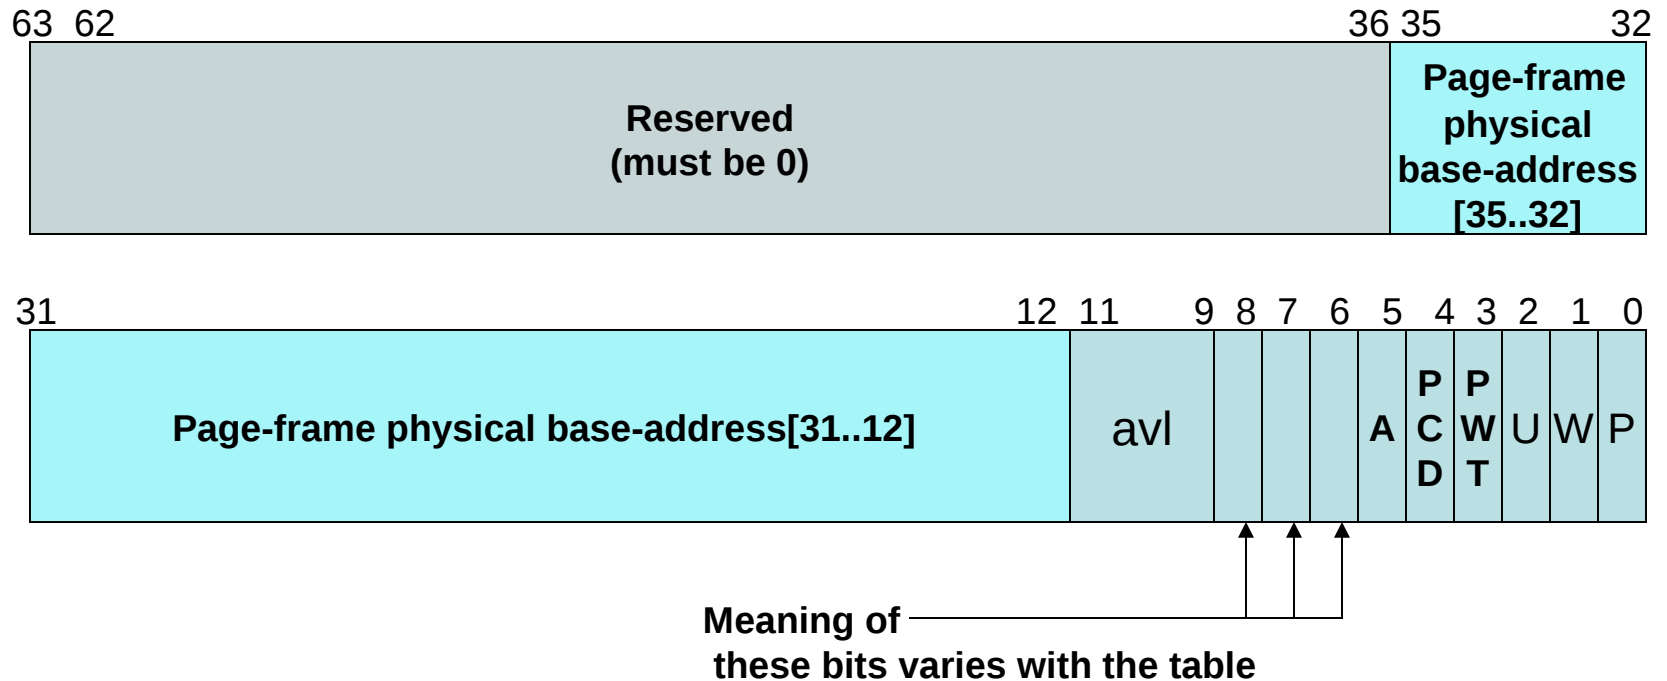
- The CPU examines any virtual address it encounters, subdividing it into **four** fields



*Each 'virtual address' is 32-bits  
but gets mapped to a physical address  
which can contain 36-bits*

# The 'pgdir' and 'pgtbl' entries

(Assumes PAE=1, but LME=0)



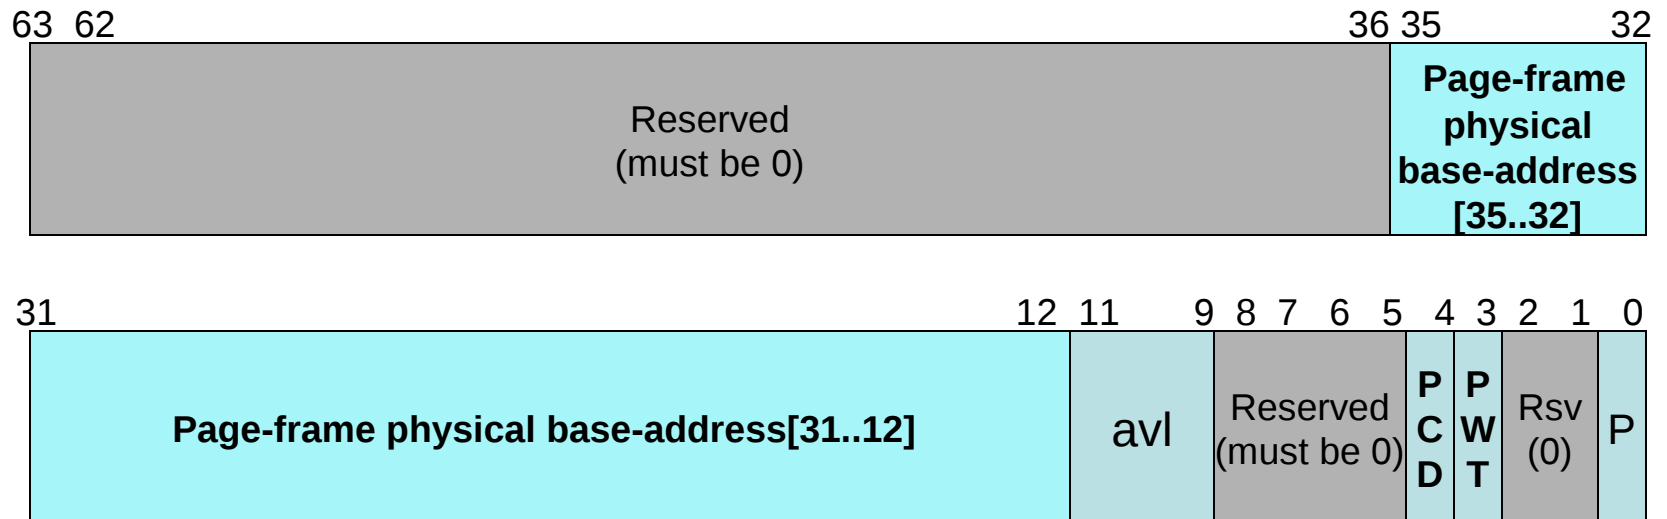
## Legend:

P = Present (1=yes, 0=no)  
W = Writable (1=yes, 0=no)  
U = User-page (1=yes, 0=no)  
A = Accessed (1=yes, 0=no)

PWT = Page Cache Disable (1=yes, 0=no)  
PWT = Page Write-Through (1=yes, 0=no)  
avl = available for user-defined purposes  
EXB = Execution-disabled Bit (if EFER.NXE=1)

# Page-directory-pointer entries

(Assumes PAE=1, but LME=0)



## Legend:

P = Present (1=yes, 0=no)

PWT = Page Cache Disable (1=yes, 0=no)

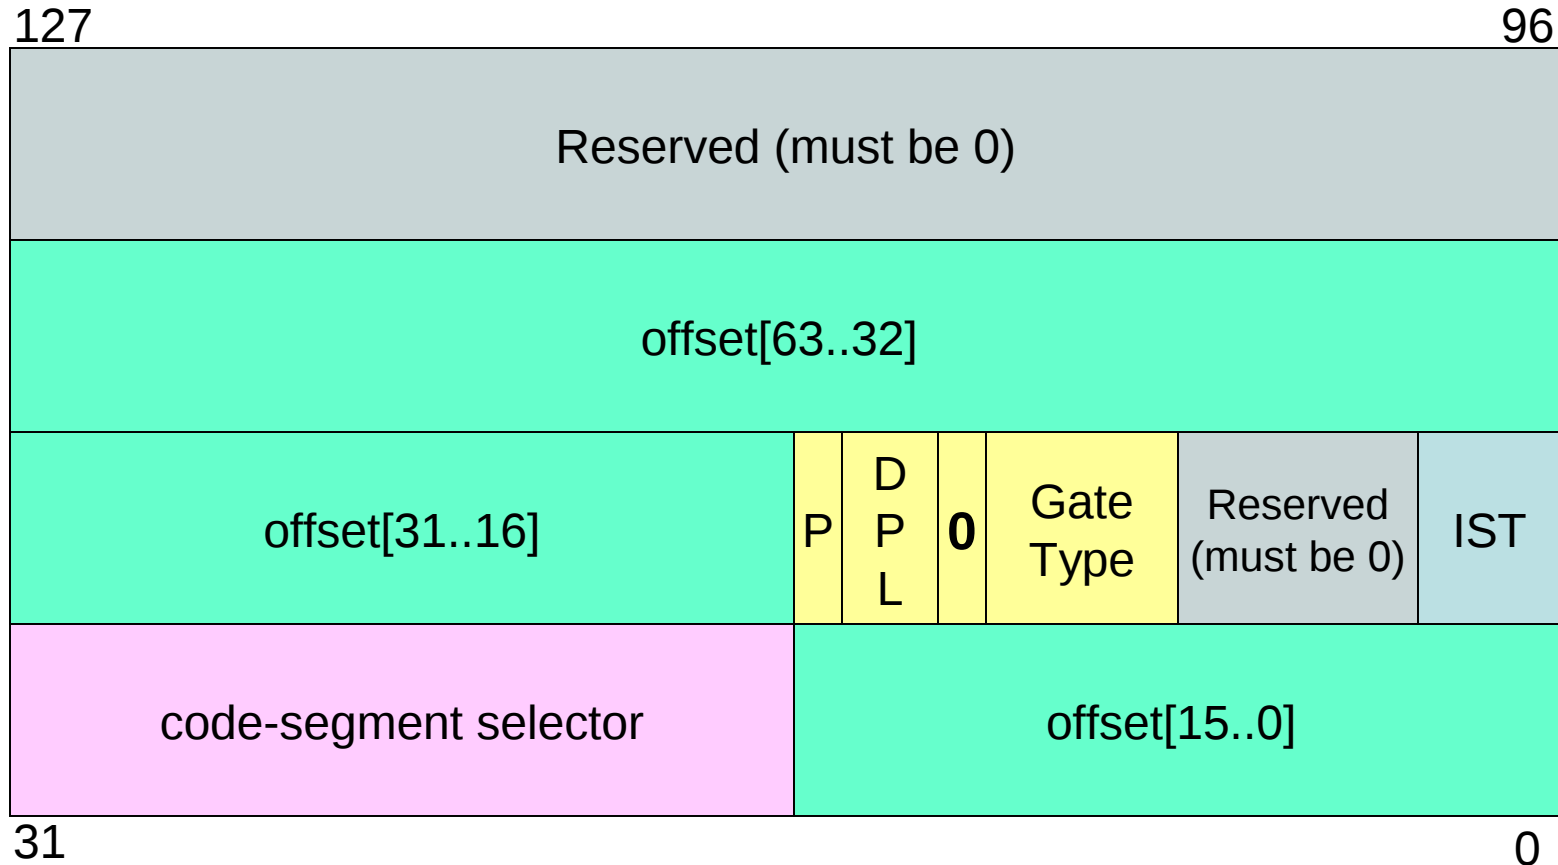
PWT = Page Write-Through (1=yes, 0=no)

avl = available for user-defined purposes

# What about IA-32e modes?

- Some differences in what amount of CPU context information is saved on the stack
- All stack-elements are quadwords (64-bits)
- All the Interrupt Service Routines execute in 64-bit code-segments (i.e., 'long mode')
- All memory-addresses involving registers CS, DS, ES, and SS are 'flat' addresses
- All the IDT gate-descriptors are 128-bits

# IA-32e IDT Gate descriptors

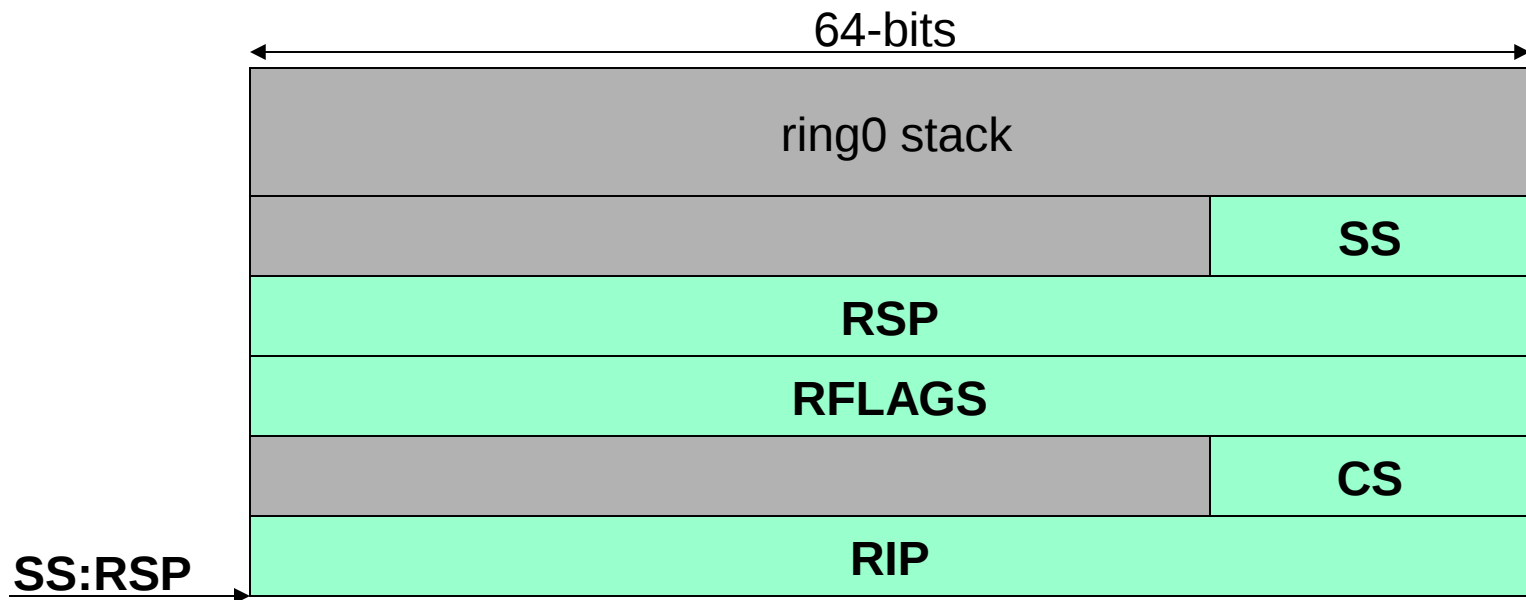


P=Present      Gate Type: 1110 = Interrupt-Gate, 1111 = Trap-Gate  
 DPL = Descriptor Privilege Level      IST = Interrupt Stack Table



# 64-bit interrupt-stackframe

- The CPU constructs a different stackframe for the 64-bit IDT gate-types – its format is independent of any privilege-changes



= always pushed

# In-class exercise #1

- Try modifying our 'tryisr32.s' program so that it does not use page-mapping at all (i.e., leave the PG-bit in CR0 turned off)
- Do you discover any problems when you try to execute that modified program?

# In-class Exercise #2

- Make a copy the 'tryisr32.s' file (named 'tryisr64.s') and see if you can discover what code-modifications are required if you wanted the timer's ISR to execute from a 64-bit code-segment (rather than from the 32-bit code-segment it formerly used); try it out on an 'anchor' machine
- (This exercise will help with Project #1)